

# **Technical Workshop - SUSE CaaS Platform v4 Beta**

## **-Workbook-**

**Course ID: Technical Workshop - SUSE Container as a Platform**

**Version: 5 Beta**

**Date: 2019-8-19**



**Proprietary Statement**

Copyright © 2018 SUSE LLC. All rights reserved.

SUSE LLC, has intellectual property rights relating to technology embodied in the product that is described in this document.

No part of this publication may be reproduced, photocopied, stored on a retrieval system, or transmitted without the express written consent of the publisher.

SUSE  
Maxfeldstrasse 5  
90409 Nuremberg  
Germany  
[www.suse.com](http://www.suse.com)

(C) 2018 SUSE LLC. All Rights Reserved. SUSE and the SUSE logo are registered trademarks of SUSE LLC in the United States and other countries. All third-party trademarks are the property of their respective owners.

If you know of illegal copying of software, contact your local Software Antipiracy Hotline.

**Disclaimer**

SUSE LLC, makes no representations or warranties with respect to the contents or use of this documentation, and specifically disclaims any express or implied warranties of merchantability or fitness for any particular purpose.

Further, SUSE LLC, reserves the right to revise this publication and to make changes to its content, at any time, without obligation to notify any person or entity of such revisions or changes. Further, SUSE LLC, makes no representations or warranties with respect to any software, and specifically disclaims any express or implied warranties of merchantability or fitness for any particular purpose. Further, SUSE LLC, reserves the right to make changes to any and all parts of SUSE software, at any time, without any obligation to notify any person or entity of such changes.

Any products or technical information provided under this Agreement may be subject to U.S. export controls and the trade laws of other countries. You agree to comply with all export control regulations and to obtain any required licenses or classification to export, re-export or import deliverables. You agree not to export or re-export to entities on the current U.S. export exclusion lists or to any embargoed or terrorist countries as specified in the U.S. export laws. You agree to not use deliverables for prohibited nuclear, missile, or chemical biological weaponry end uses. SUSE assumes no responsibility for your failure to obtain any necessary export approvals.

This SUSE Training Manual is published solely to instruct students in the use of SUSE networking software. Although third-party application software packages may be used in SUSE training courses, this is for demonstration purposes only and shall not constitute an endorsement of any of these software applications.

Further, SUSE LLC does not represent itself as having any particular expertise in these application software packages and any use by students of the same shall be done at the student's own risk.

# Table of Contents

Documentation Conventions:.....	6
<b>Section 1 : Demonstration of SUSE CaaS Platform.....</b>	<b>9</b>
(No Exercises).....	10
Exercise 1 : Start the Lab VMs.....	11
Task 1: Start the Lab VMs.....	11
Exercise 2 : Install a CaaS Platform Nodes.....	13
Task 1: Deploy the Kubernetes Dashboard.....	13
Task 2: Test to ensure SSH keys are properly set.....	14
Exercise 3 : Setup SSH Keys.....	15
Task 1: Deploy the Kubernetes Dashboard.....	15
Task 2: Run ssh-keygen on all nodes.....	15
Task 3: Copy SSH keys to Management workstation.....	16
Task 4: Test to ensure SSH keys are properly set.....	16
Exercise 4 : Install a CaaS Platform Nodes.....	18
Task 1: Deploy the first Kubernetes Master node.....	18
<b>Section 2 : Connecting to the Cluster.....</b>	<b>21</b>
Exercise 1 : Use the kubectl Command to Display Info About the Cluster.....	21
Task 1: Configure workstation for kubectl command.....	21
Task 2: Use the kubectl Command.....	21
Exercise 2 : Deploy the Kubernetes Dashboard.....	23
Task 1: Deploy the Kubernetes Dashboard.....	24
Task 2: Access the Kubernetes Dashboard.....	24
Exercise 3 : Install and Configure Helm.....	25
Task 1: Install Helm on Workstation.....	25
Task 2: Create Service account and Cluster Role for Helm.....	26
Task 3: Initialize Helm.....	26
Task 4: Add SUSE Charts to Helm.....	27
Task 5: Install Centralize Logging.....	28
<b>Section 3 : Deploying a Workload.....</b>	<b>28</b>
Exercise 1 : Deploy a Simple Pod on Kubernetes.....	28
Task 1: View a Manifest for the Deployment.....	28
Task 2: Deploy a simple application (GUI option).....	30
Task 3: Deploy a simple application (Command Line option).....	32
Exercise 2 : Delete a Deployment on Kubernetes.....	33
Task 1: Delete the Deployment (GUI Option).....	33
Task 2: Delete the Deployment (Command Line option).....	34
Exercise 3 : Scale Out a Deployment.....	35
Task 1: View the Manifest for the Deployment.....	35
Task 2: Scale the Deployment (GUI option).....	35
Task 3: Scale the Deployment (Command Line Option).....	38
<b>Section 4 : Work With Kubernetes.....</b>	<b>39</b>

Exercise 1 : Update a Deployment.....	39
Task 1: Create a New Manifest for the Deployment.....	39
Task 2: Update the Deployment.....	40
Exercise 2 : Update a Deployment Via Rolling Updates.....	41
Task 1: Create a New Manifest for the Deployment.....	41
Task 2: Update the Deployment.....	42
Exercise 3 : Expose a Service Running in a Pod.....	42
Task 1: Create a Manifest for the Service.....	43
Task 2: Define the Service.....	43
Task 3: Access the Exposed Service.....	43
Exercise 4 : Setup Readiness and Liveness Probes.....	44
Task 1: View manifest for the Deployment.....	44
Task 2: Deploy Manifest.....	46
Task 3: View the Replica sets.....	46
Task 4: Kill an nginx server.....	46
Exercise 5 : Define Limits for Containers and Pods in Kubernetes.....	47
Task 1: Create a New Namespace in the Cluster.....	48
Task 2: Create a Manifest for the Limits.....	48
Task 3: Apply the Limits to the Namespace.....	48
Exercise 6 : Introduction to Helm.....	49
Task 1: Use Some Basic Helm Commands.....	49
Exercise 7 : Deploy an Application with Helm.....	50
Task 1: Create Helm Chart Config File.....	50
Task 2: Deploy the Helm Chart.....	51
Task 3: Access the Application Deployed by the Helm Chart.....	51
Task 4: Update the Application Release.....	52
Task 5: Delete a Deployed Helm Chart Release.....	53

## **Section 5 : Deploying Advanced Workload.....53**

Exercise 1 : Run Heimdall as a standalone container.....	53
Task 1: Create a place for Heimdall to store it's files.....	54
Task 2: Run Heimdall as a container.....	54
Task 3: On the Workstation launch the Chrome Browser.....	55
Task 4: Create an App in Helm.....	56
Task 5: Stop and restart container.....	57
Task 6: Verify with Browser.....	57
Task 7: Stop Container Instance.....	57
Task 8: Try and run multiple versions of Heimdall.....	57
Exercise 2 : Launch Heimdall.....	58
Task 1: Create the Manifest for the Persistent Volume.....	58
Task 2: Deploy Heimdall.....	59
Task 3: View the Deployment.....	60
Task 4: View Heimdall in a Browser.....	62
Task 4: Re-deploy Heimdall.....	62
Exercise 3 : Configure Heimdall to use NFS Persistent Storage.....	62
Task 1: Create the Manifest for the Persistent Volume.....	62
Task 2: Create the Manifest for the Persistent Volume Claim.....	63

SUSE U Advanced - SUSE CaaS Platform

Task 3: Create the Manifest for the Service so we can access the deployment.....	63
Task 4: Create the Manifests for the Web Frontend.....	64
Task 5: Deploy the Objects.....	65
Task 6: Test the Persistent Data.....	66
Task 7: Remove the Objects from the Cluster.....	66

## Documentation Conventions:

---

The following typographical conventions are used in this manual:

<b>Bold</b>	Represents things you should pay attention to or buttons you click, text or options that you should click/select/type in a GUI.
<b>Bold Gray</b>	Represents the name of a Task or in the context of what is seen on the screen, the screen name, a tab name, column name, field name, etc.
<b>Bold Red</b>	Represents warnings or very important information.
<b>Option &gt; Option &gt; Option</b>	Represents a chain of items selected from a menu.
<b><i>BOLD_UPPERCASE_ITALIC</i></b>	Represents an “exercise variable” that you replace with another value.
<b>bold monospace</b>	Represents text displayed in a terminal or entered in a file.
<b>bold monospace blue</b>	Represents commands entered at the command line.
<b>bold monospace green</b>	Represents a file name.







# 1 Demonstration of SUSE CaaS Platform

---

**Description:**

How to highlight the Key Benefits of SUSE CaaS Platform

**(No Exercises)**

# 1- 1 Start the Lab VMs

## Description:

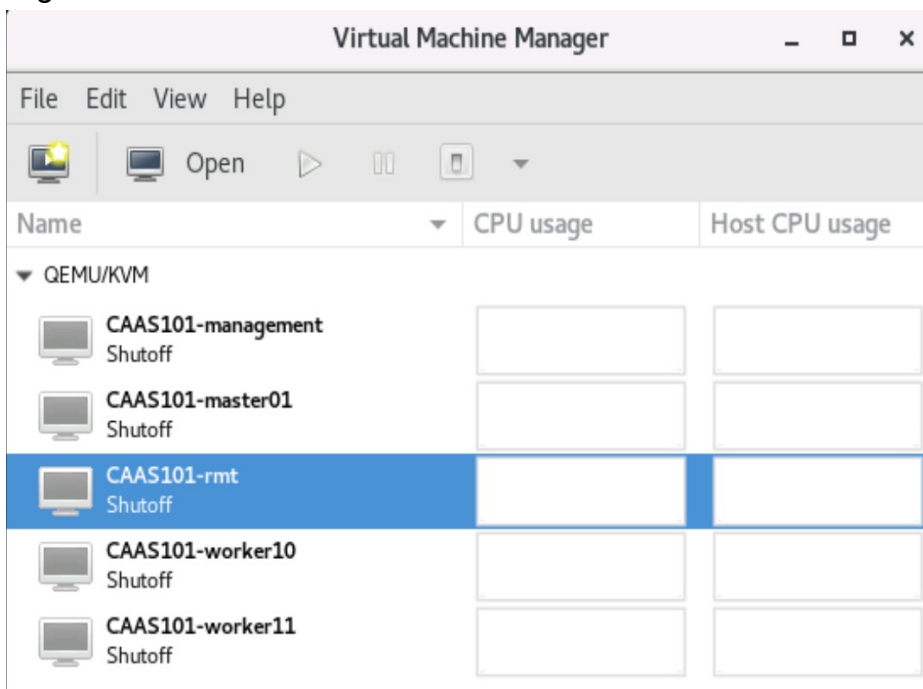
In this exercise, you use the Virt-Manager utility to start the lab VM(s) in the proper order.

## Task 1: Start the Lab VMs

1. On the **VM host**, launch the Virt-Manager utility  
(4<sup>th</sup> icon from the left on the dock at the bottom of the screen)  
You should see the course VMs listed.



2. Right-click on the **CAAS101-rmt** VM and select **Run**




3. Double-click on the **CAAS101-rmt** VM to view its console  
When the **CAAS101-rmt** VM has finished booting completely, close its console window.  
(You can use the CPU utilization graph in Virt-Manager to see when the node is finished booting and starting services).

4. Power on the reset of the Virtual Machines **CAAS101-management**, **CAAS101-master**, **CAAS101-worker10** and **CAAS101-worker10**

\*note you can power all of the on at the same time

When the **Workstation** has finished booting you will use it as your management workstation for all tasks including managing VMs with the Virt-Manager utility.

5. Make sure you double click on the **CAAS101-management** so you view it full screen
6. Open a Terminal session by clicking on the 3<sup>rd</sup> icon down



---

## Summary:

In this exercise, you started the required lab VM(s) in the proper order.

(End of Exercise)

# 1-2 Install a CaaS Platform Nodes

## Description:

In this exercise, you install a SUSE CaaS Platform all nodes.

## Dependencies:

The RMT server must be configured (Optional Docker Registry)

DHCP, PXE and DNS must be configured.

Master, and 2 worker nodes with SLES 15 SP1 installed and CaaS Platform Repos added

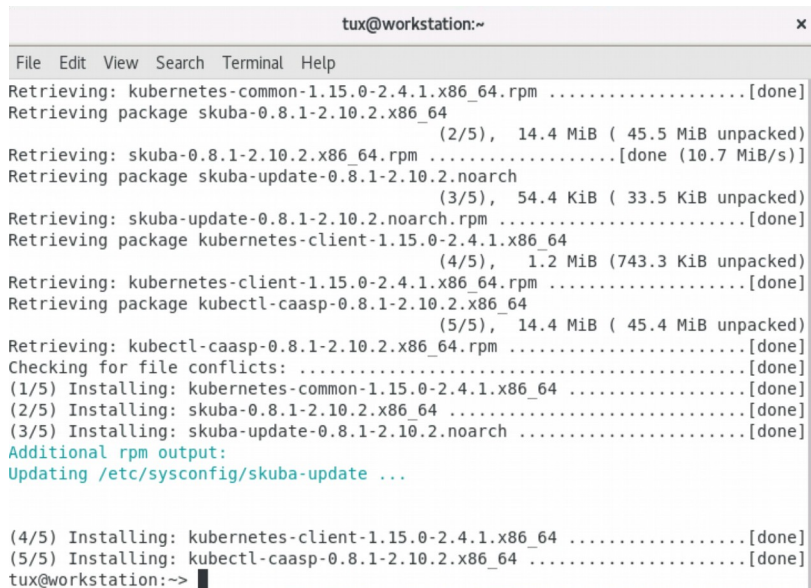
## Task 1: Deploy the Kubernetes Dashboard

1. On the **Workstation**, open a terminal
2. Enter the following command to install skuba and the kubernetes client

**sudo zypper in skuba kubernetes-client**

I will tell you the amount of additional drive space needed and it will prompt you to press **y** to continue

3. Press the '**Y**' key to continue  
when it finished you should see a screen similar to the image below



```
tux@workstation:~
File Edit View Search Terminal Help
Retrieving: kubernetes-common-1.15.0-2.4.1.x86_64.rpm .....[done]
Retrieving package skuba-0.8.1-2.10.2.x86_64
(2/5), 14.4 MiB ( 45.5 MiB unpacked)
Retrieving: skuba-0.8.1-2.10.2.x86_64.rpm .....[done (10.7 MiB/s)]
Retrieving package skuba-update-0.8.1-2.10.2.noarch
(3/5), 54.4 KiB ( 33.5 KiB unpacked)
Retrieving: skuba-update-0.8.1-2.10.2.noarch.rpm .....[done]
Retrieving package kubernetes-client-1.15.0-2.4.1.x86_64
(4/5), 1.2 MiB (743.3 KiB unpacked)
Retrieving: kubernetes-client-1.15.0-2.4.1.x86_64.rpm .....[done]
Retrieving package kubectl-caasp-0.8.1-2.10.2.x86_64
(5/5), 14.4 MiB ( 45.4 MiB unpacked)
Retrieving: kubectl-caasp-0.8.1-2.10.2.x86_64.rpm .....[done]
Checking for file conflicts: .....[done]
(1/5) Installing: kubernetes-common-1.15.0-2.4.1.x86_64 .....[done]
(2/5) Installing: skuba-0.8.1-2.10.2.x86_64 .....[done]
(3/5) Installing: skuba-update-0.8.1-2.10.2.noarch .....[done]
Additional rpm output:
Updating /etc/sysconfig/skuba-update ...

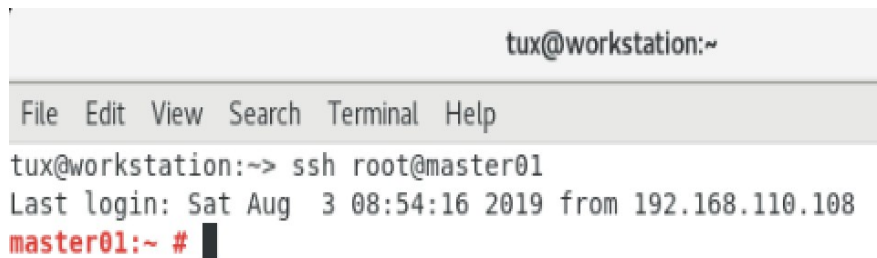
(4/5) Installing: kubernetes-client-1.15.0-2.4.1.x86_64 .....[done]
(5/5) Installing: kubectl-caasp-0.8.1-2.10.2.x86_64 .....[done]
tux@workstation:~>
```

## Task 2: Test to ensure SSH keys are properly set

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to test your **ssh keys** to master01

**ssh root@master01**

3. You should see a red prompt that says master01



```
tux@workstation:~  
File Edit View Search Terminal Help  
tux@workstation:~> ssh root@master01  
Last login: Sat Aug 3 08:54:16 2019 from 192.168.110.108  
master01:~ #
```

4. Notice that you did not need to type in a password or accept a key?  
This is because of the pre-work done on the workstation
5. Exit your ssh session from the terminal prompt  
**exit**
6. Repeat the step 2 and 3 for **worker10** and **worker11**

## Summary:

In this exercise, you installed Skuba and the Kubernetes Client Utilities. You also tested that the SSH keys were set up properly and you could seamlessly connect to all of the nodes.

(End of Exercise)

## 1- 3 Setup SSH Keys

---

### Description:

In this exercise, you install setup all of the machines so you can ssh to any of the nodes with a username or password

---

### Task 1: Deploy the Kubernetes Dashboard

1. On the **Workstation**, open a terminal
2. Enter the following command to Start the ssh-agent  
**eval "\$(ssh-agent -s)"**
3. Enter the following command to create an ssh key  
**ssh-keygen**  
When it asks you questions just take the default

### Task 2: Run ssh-keygen on all nodes

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to test your **ssh keys** to master01  
**ssh root@master01**  
password is **linux**
3. You should see a red prompt that says master01

```
tux@workstation:~
File Edit View Search Terminal Help
tux@workstation:~> ssh root@master01
Last login: Sat Aug 3 08:54:16 2019 from 192.168.110.108
master01:~ #
```

4. Enter the following command in the terminal to create an **ssh key**  
**ssh-keygen**
5. Exit your ssh session from the terminal prompt  
**exit**
6. Repeat the step 2 and 3 for **worker10** and **worker11**

### Task 3: Copy SSH keys to Management workstation

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to test your **ssh keys** to master01  
**ssh-copy-id -i ~/.ssh/id\_rsa.pub root@master**
3. Repeat the step 2 and 3 for **worker10** and **worker11**

### Task 4: Test to ensure SSH keys are properly set

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to test your **ssh keys** to master01  
**ssh root@master01**
3. You should see a red prompt that says master01

```
tux@workstation:~
File Edit View Search Terminal Help
tux@workstation:~> ssh root@master01
Last login: Sat Aug 3 08:54:16 2019 from 192.168.110.108
master01:~ #
```

4. Notice that you did not need to type in a password or accept a key?  
This is because of the pre-work done on the workstation
5. Exit your ssh session from the terminal prompt



**exit**

6. Repeat the step 2 and 3 for **worker10** and **worker11**

## Summary:

In this exercise, you set it up so that you can ssh from the Management Workstation to any of the nodes without asking for a username or a password

(End of Exercise)

## 1- 4 Install a CaaS Platform Nodes

---

### Description:

In this exercise, you install a SUSE CaaS Platform all nodes.

#### Dependencies:

The RMT server must be configured (Optional Docker Registry)

DHCP, PXE and DNS must be configured.

Master, and 2 worker nodes with SLES 15 SP1 installed and CaaS Platform Repos added

---

### Task 1: Deploy the first Kubernetes Master node

1. On the **Workstation**, open a terminal
2. Enter the following command create the basic config to our **kubernetes cluster**

**skuba cluster init --control-plane master.example.com my-cluster**

```
tux@workstation:~> skuba cluster init --control-plane master.example.com my-cluster
** This is a BETA release and NOT intended for production usage. **
[init] configuration files written to /home/tux/my-cluster
tux@workstation:~> █
```

3. Enter the following command to change into the newly created folder

**cd my-cluster**

4. Look at the cluster config filename **my-cluster****less kubeadm-init.conf**

```

tux@workstation:~/my-cluster> cat kubeadm-init.conf
apiVersion: kubeadm.k8s.io/v1beta1
kind: InitConfiguration
bootstrapTokens: []
localAPIEndpoint:
  advertiseAddress: ""
---
apiVersion: kubeadm.k8s.io/v1beta1
kind: ClusterConfiguration
apiServer:
  certSANs:
    - master.example.com
  extraArgs:
    oidc-issuer-url: https://master.example.com:32000
    oidc-client-id: oidc
    oidc-ca-file: /etc/kubernetes/pki/ca.crt
    oidc-username-claim: email
    oidc-groups-claim: groups
clusterName: my-cluster
controlPlaneEndpoint: master.example.com:6443
dns:
  imageRepository: registry.suse.com/caasp/v4
  imageTag: 1.3.1
  type: CoreDNS
etcd:
  local:
    imageRepository: registry.suse.com/caasp/v4
    imageTag: 3.3.11
imageRepository: registry.suse.com/caasp/v4
kubernetesVersion: 1.15.0
networking:
  podSubnet: 10.244.0.0/16
  serviceSubnet: 10.96.0.0/12
useHyperKubeImage: true
tux@workstation:~/my-cluster> █

```

## 5. Install, configure and bootstrap the cluster-info

**skuba node bootstrap --target master.example.com master01**

```
tux@workstation:~/my-cluster> skuba node bootstrap --target master.example.com master01
** This is a BETA release and NOT intended for production usage. **
W0808 17:13:25.590213 2588 ssh.go:306]
The authenticity of host '192.168.110.7:22' can't be established.
ECDSA key fingerprint is 54:3e:9b:2e:21:fc:5c:70:ff:c4:99:7c:fa:b4:08:d0.
I0808 17:13:25.590316 2588 ssh.go:307] accepting SSH key for "master.example.com:22"
I0808 17:13:25.590327 2588 ssh.go:308] adding fingerprint for "master.example.com:22" to "known_hosts"
[bootstrap] updating init configuration with target information
[bootstrap] writing init configuration for node
[bootstrap] applying init configuration to node
[bootstrap] downloading secrets from bootstrapped node "master.example.com"
[bootstrap] successfully bootstrapped node "master.example.com" with Kubernetes: "1.15.0"
tux@workstation:~/my-cluster>
```

6. add a **worker10** node to the cluster

**skuba node join --role worker --target worker10.example.com worker10**

```
tux@workstation:~/my-cluster> skuba node join --role worker --target worker10.example.com worker10
** This is a BETA release and NOT intended for production usage. **
W0808 17:23:50.580677 2632 ssh.go:306]
The authenticity of host '192.168.110.10:22' can't be established.
ECDSA key fingerprint is 19:7c:db:38:d2:c5:34:f6:cb:90:21:4e:1a:b5:5e:d7.
I0808 17:23:50.580863 2632 ssh.go:307] accepting SSH key for "worker10.example.com:22"
I0808 17:23:50.580916 2632 ssh.go:308] adding fingerprint for "worker10.example.com:22" to "known_hosts"
[join] applying states to new node
[join] node successfully joined the cluster
tux@workstation:~/my-cluster>
```

7. add a **worker11** node to the cluster

**skuba node join --role worker --target worker11.example.com worker11**

You should see a message very similar to the message you received in the previous step.

## 8. Verify Cluster Status

**skuba cluster status**

```
tux@workstation:~/my-cluster> skuba cluster status
** This is a BETA release and NOT intended for production usage. **


| NAME     | OS-IMAGE                            | KERNEL-VERSION         | KUBELET-VERSION | CONTAINER-RUNTIME | HAS-UPDATES | HAS-DISRUPTIVE-UPDATES |
|----------|-------------------------------------|------------------------|-----------------|-------------------|-------------|------------------------|
| master01 | SUSE Linux Enterprise Server 15 SP1 | 4.12.14-197.10-default | v1.15.0         | cri-o://1.15.0    | <none>      | <none>                 |
| worker10 | SUSE Linux Enterprise Server 15 SP1 | 4.12.14-197.10-default | v1.15.0         | cri-o://1.15.0    | <none>      | <none>                 |
| worker11 | SUSE Linux Enterprise Server 15 SP1 | 4.12.14-197.10-default | v1.15.0         | cri-o://1.15.0    | <none>      | <none>                 |


tux@workstation:~/my-cluster>
```

## Summary:

In this exercise, you installed Skuba and the Kubernetes Client Utilities. You also tested that the SSH keys were set up properly and you could seamlessly connect to all of the nodes.

(End of Exercise)

## 2 Connecting to the Cluster

---

### Description:

In this section you will learn how to connect to the Cluster via both the command line and a GUI interface.

### 2- 1 Use the kubectl Command to Display Info About the Cluster

---

### Description:

In this exercise, you use the **kubectl** command to display info about the Kubernetes cluster.

---

### Task 1: Configure workstation for kubectl command

1. On the **Workstation**, open a terminal
2. Create official folder for kubernetes config filename

```
mkdir ~/.kube
```

3. On the **Workstation**, open a terminal
4. Create config file

```
cp ~/my-cluster/admin.conf ~/.kube/config
```

### Task 2: Use the kubectl Command

1. From a terminal, enter the following command to view the URL of the Kubernetes master:

```
kubectl cluster-info
```

You should see the URL of the Kubernetes master displayed.

2. Enter the following command to see a more detailed output of the status of the cluster:

```
kubectl cluster-info dump
```

You should see a json dump of the status of cluster.

3. Enter the following command to display a list of Kubernetes nodes:

### **kubectl get nodes**

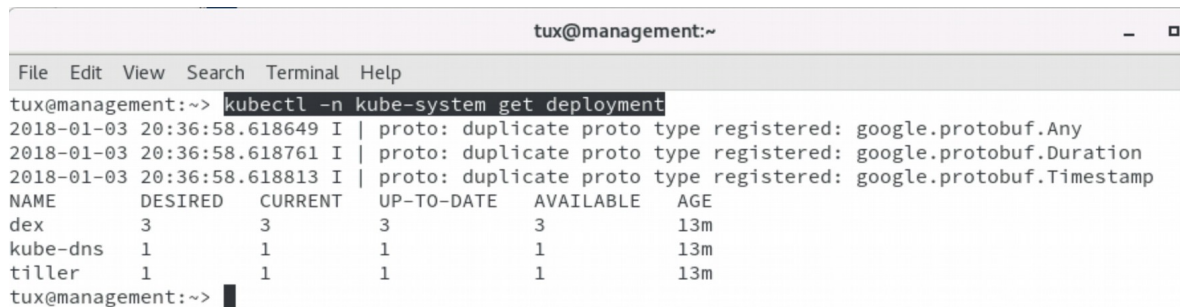
You should see the list of nodes, their node names, status, age and version displayed.

```
tux@management:~/course_files/CAAS101/scripts> kubectl get nodes
2018-08-08 10:02:50.924418 I | proto: duplicate proto type registered: google.protobuf.Any
2018-08-08 10:02:50.924462 I | proto: duplicate proto type registered: google.protobuf.Duration
2018-08-08 10:02:50.924473 I | proto: duplicate proto type registered: google.protobuf.Timestamp
NAME          STATUS    AGE           VERSION
master01      Ready     1d            v1.9.8
master02      Ready     1d            v1.9.8
master03      Ready     1d            v1.9.8
worker10      Ready     1d            v1.9.8
worker11      Ready     1d            v1.9.8
worker12      Ready     1d            v1.9.8
tux@management:~/course_files/CAAS101/scripts> █
```

4. Enter the following command to display a list of Kubernetes nodes:

### **kubectl -n kube-system get deployments**

You should see the list of deployments, their nodes names, status, age and version displayed.



```
tux@management:~> kubectl -n kube-system get deployment
2018-01-03 20:36:58.618649 I | proto: duplicate proto type registered: google.protobuf.Any
2018-01-03 20:36:58.618761 I | proto: duplicate proto type registered: google.protobuf.Duration
2018-01-03 20:36:58.618813 I | proto: duplicate proto type registered: google.protobuf.Timestamp
NAME          DESIRED    CURRENT    UP-TO-DATE    AVAILABLE    AGE
dex           3          3          3             3            13m
kube-dns      1          1          1             1            13m
tiller        1          1          1             1            13m
tux@management:~> █
```

5. Enter the following command to display a list of Kubernetes pods:

**kubectl -n kube-system get pods**

You should see the list of pods, their names, status, and age displayed.

```
tux@management:~
File Edit View Search Terminal Help
tux@management:~> kubectl -n kube-system get pods
2018-01-03 20:34:37.713518 I | proto: duplicate proto type registered: google.protobuf.Any
2018-01-03 20:34:37.713579 I | proto: duplicate proto type registered: google.protobuf.Duration
2018-01-03 20:34:37.713607 I | proto: duplicate proto type registered: google.protobuf.Timestamp
NAME                                READY    STATUS    RESTARTS   AGE
dex-3532614258-cssdz                1/1      Running   2           10m
dex-3532614258-nzjvq                1/1      Running   2           10m
dex-3532614258-xvpc9                1/1      Running   2           10m
haproxy-6c4cbefa4e0545fc97a891c78b2f2b1c.infra.caasp.local 1/1      Running   0           9m
haproxy-b099b425f0194d18902aa68729c2a12e.infra.caasp.local 1/1      Running   0           10m
haproxy-d115a10918394ac88ab8c88bf7350556.infra.caasp.local 1/1      Running   0           10m
haproxy-dda6d06e26214c6d9c208ec2899748d2.infra.caasp.local 1/1      Running   0           10m
haproxy-e8185390bbc544efba4965e117a960eb.infra.caasp.local 1/1      Running   0           10m
haproxy-f7f84cefb02e4eadbc03ce48642eef.infra.caasp.local 1/1      Running   0           9m
kube-dns-1144198277-3zf34           3/3      Running   0           11m
tiller-3881891813-4251j             1/1      Running   0           10m
tux@management:~>
```

## Summary:

In this exercise, you downloaded the kubectl config file from the cluster and then ran some kubectl commands to view information about the cluster.

(End of Exercise)

## 2- 2 Deploy the Kubernetes Dashboard

### Description:

In this exercise, you deploy the Kubernetes Dashboard.

We will also apply an RBAC configuration that turns off token based authentication and stops the dashboard from timing out after 10 mins.

## Task 1: Deploy the Kubernetes Dashboard

1. On the **Workstation**, open a terminal
2. Enter the following command to deploy the Kubernetes Dashboard:

```
kubectl apply -f ~/STW-CaaSPv4/manifests/dashboard
```

You should see that a **secret**, **serviceaccount**, **role**, **rolebinding**, **deployment** and **service** were created:

3. Enter the following command to view the pods deployed in the **kube-system** namespace:

```
kubectl -n kube-system get pods
```

You should see the kubernetes-dashboard listed among the pods.

4. Enter the following command to view the deployments in the **kube-system** namespace:

```
kubectl -n kube-system get deployments
```

You should see the kubernetes-dashboard listed among the deployments.

5. Enter the following command to view the services in the **kube-system** namespace:

```
kubectl -n kube-system get services
```

You should see the kubernetes-dashboard listed among the services.

## Task 2: Access the Kubernetes Dashboard

1. Enter the following command to start a kubectl proxy in a screen session:

```
screen kubectl proxy
```

You should see the proxy session started. Notice that port 8001 on the local host is the port used.

2. Enter the following keystrokes to detach from the screen session:

```
ctrl+a ctrl+d
```

You should be detached from the screen session and back at a command prompt.

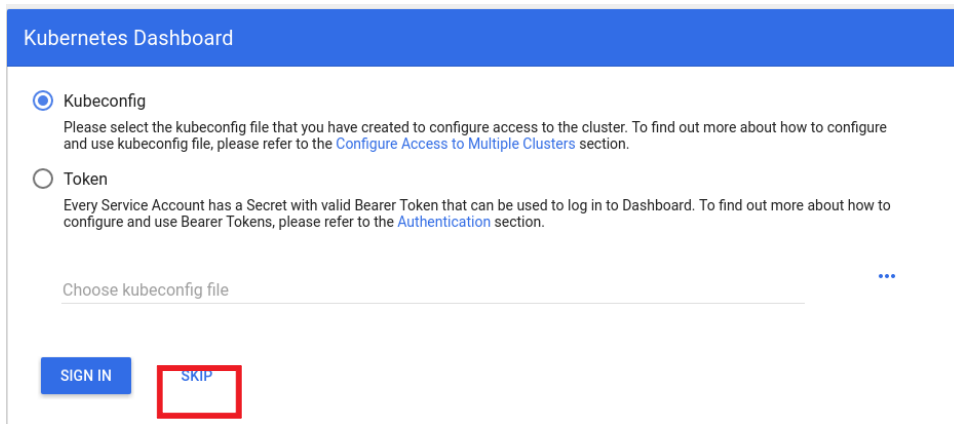
3. Open a web browser and point to:

```
http://localhost:8001/api/v1/namespaces/kube-system/services/  
https:kubernetes-dashboard:/proxy/#!/login
```

You should see the Kubernetes Dashboard authentication method screen.

4. Click: **Skip** to login without a token





5. From the list on the left, in the **Namespace** section, from the drop-down list (which is probably displaying the **default** namespace), select **kube-system**  
You should see the kubernetes dashboard deployment in the **kube-system** namespace.
6. From that same drop-down menu, select **default**  
You should see that there is nothing deployed in the **default** namespace.

---

### Summary:

In this exercise, you deployed the Kubernetes Dashboard to the cluster. You then started a kubectl proxy in a screen session and accessed the Kubernetes Dashboard.

(End of Exercise)

## 2- 3 Install and Configure Helm

---

### Description:

In this exercise, you install install Helm and config it to pull charts from SUSE.

---

### Task 1: Install Helm on Workstation

1. On the **Workstation**, open a terminal
2. Enter the following command to install skuba and the kubernetes client

**sudo zypper in helm**

I will tell you the amount of additional drive space needed and it will prompt you to press

y to continue

## Task 2: Create Service account and Cluster Role for Helm

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to create a service account for Helm to use  
**kubectl create serviceaccount --namespace kube-system tiller**
3. Enter the following command in the terminal to create a cluster role for Helm to use

**kubectl create clusterrolebinding tiller --clusterrole=cluster-admin --serviceaccount=kube-system:tiller**

## Task 3: Initialize Helm

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to create a service account for Helm to use  
**helm init --tiller-image registry.suse.com/caasp/v4/helm-tiller:2.8.2 --service-account tiller**

- Look to see if tiller is currently running (give it a minute if it is not in the **Ready** state)

### kubectl get pods --all-namespaces

```
tux@localhost:~> kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	cilium-82kz9	1/1	Running	0	47m
kube-system	cilium-bzsp7	1/1	Running	0	47m
kube-system	cilium-operator-7d6dddbf5-qnpb5	1/1	Running	0	50m
kube-system	cilium-rwwjw	1/1	Running	0	47m
kube-system	coredns-69c4947958-75vdp	1/1	Running	1	50m
kube-system	coredns-69c4947958-vmf6d	1/1	Running	3	50m
kube-system	etcd-master01	1/1	Running	0	49m
kube-system	kube-apiserver-master01	1/1	Running	0	49m
kube-system	kube-controller-manager-master01	1/1	Running	0	49m
kube-system	kube-proxy-2nhh6	1/1	Running	0	50m
kube-system	kube-proxy-ftbck	1/1	Running	0	47m
kube-system	kube-proxy-pj698	1/1	Running	0	48m
kube-system	kube-scheduler-master01	1/1	Running	0	49m
kube-system	kubernetes-dashboard-5f9c6b756f-hkwxh	1/1	Running	0	45m
kube-system	kured-4qzsf	1/1	Running	0	45m
kube-system	kured-7tcfj	1/1	Running	0	49m
kube-system	kured-9z4g7	1/1	Running	0	47m
kube-system	oidc-dex-55fc689dc-mccw8	1/1	Running	1	50m
kube-system	oidc-gangway-7b7fbdbdf-4w6z8	1/1	Running	0	50m
kube-system	tiller-deploy-7c666b7c99-7whpp	1/1	Running	0	42m

```
tux@localhost:~> █
```

## Task 4: Add SUSE Charts to Helm

- On the **Workstation**, open a terminal
- Enter the following command in the terminal to create a service account for Helm to use  
**helm repo add suse-charts <https://kubernetes-charts.suse.com>**
- Look at available charts from SUSE

### helm search suse

```
tux@localhost:~> helm search suse
```

NAME	CHART VERSION	APP VERSION	DESCRIPTION
suse-charts/cf	2.17.1	1.4.1	A Helm chart for SUSE Cloud Foundry
suse-charts/cf-usb-sidecar-mysql	1.0.1		A Helm chart for SUSE Universal Service Broker ...
suse-charts/cf-usb-sidecar-postgres	1.0.1		A Helm chart for SUSE Universal Service Broker ...
suse-charts/console	2.4.0	2.4.0	A Helm chart for deploying Stratos UI Console
suse-charts/log-agent-rsyslog	1.0.1	8.39.0	Log Agent for forwarding logs of K8s control pl...
suse-charts/metrics	1.0.0	1.0.0	A Helm chart for Stratos Metrics
suse-charts/minibroker	0.2.0		A minibroker for your minikube
suse-charts/nginx-ingress	0.28.4	0.15.0	An nginx Ingress controller that uses ConfigMap...
suse-charts/uaa	2.17.1	1.4.1	A Helm chart for SUSE UAA

```
tux@localhost:~> █
```

## Task 5: Install Centralize Logging

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to create a service account for Helm to use

```
helm install suse-charts/log-agent-rsyslog --name 1.0.1 --set server.host=rsyslog-server.default.svc.cluster.local --set server.port=514
```

---

### Summary:

In this exercise, you installed Helm on the Management workstation. You then setup and configured it to pull Charts directly from SUSE. You then installed the CaaS Platform's Centralized Logging Service.

(End of Exercise)

## 3 Deploying a Workload

---

### Description:

In this section you will deploy a simple workload in Kubernetes.

### 3- 1 Deploy a Simple Pod on Kubernetes

---

### Description:

In this exercise, you deploy the Nginx web server as a simple pod on the Kubernetes cluster.

---

### Task 1: View a Manifest for the Deployment

1. On the **management workstation**, in the text editor of your choice, open the file:

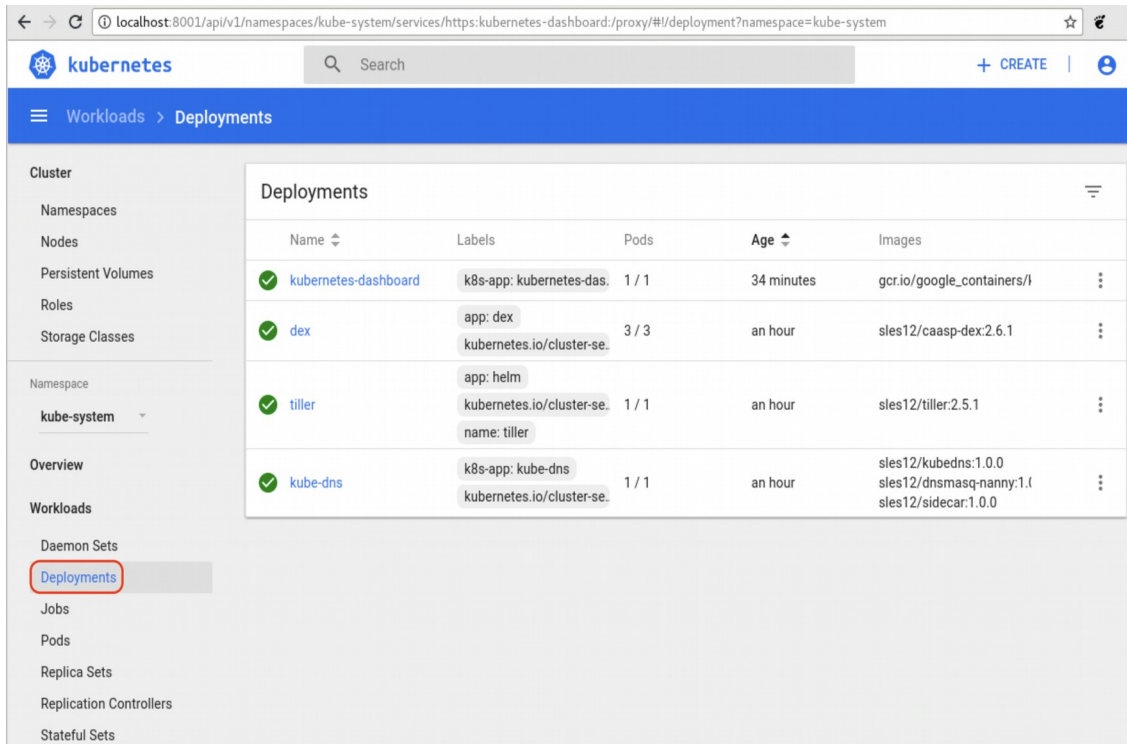
```
~/STW-CaaS-Pv4/labs/nginx-deployment.yaml
```

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 4
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: smt.example.com:5000/sles12sp3_nginx
        ports:
        - containerPort: 80
```

## Task 2: Deploy a simple application (GUI option)

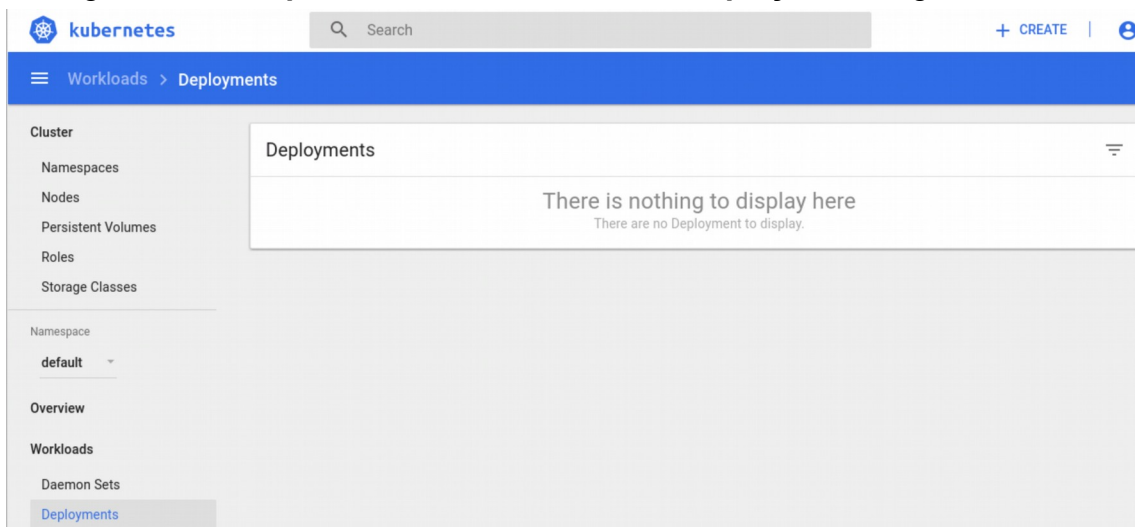
To deploy the application, go to the Kubernetes Dashboard

1. Click on **Deployments** on the left hand panel



All of the currently deployed Deployments are in the kube-system namespace

2. Select **Namespace** on the left hand panel
3. Change the **Namespace** to **Default** and select **Deployments** again



note there are no current deployments in the **default** namespace

4. Click on **+ CREATE**
5. Select **Create from File** and then press '...' to select file

CREATE FROM TEXT INPUT

CREATE FROM FILE

CREATE AN APP

Select YAML or JSON file specifying the resources to deploy to the currently selected namespace. [Learn more](#)

Choose YAML or JSON file

UPLOAD

CANCEL

6. Browse to **~/STW-CaaSv4/labs/nginx-deployment.yaml**
7. Press Upload
8. Select **Deployments** view and we should now see the nginx-deployment

kubernetes

Search

+ CREATE |

Workloads > Deployments

Cluster

Namespaces

Nodes

Persistent Volumes

Roles

Storage Classes

Namespace

default

Overview

Workloads

Daemon Sets

Deployments

Deployments

Name	Labels	Pods	Age	Images
✓ nginx-deployment	app: nginx	1 / 1	32 seconds	nginx:1.7.9

## 9. Explore the Kubernetes Dashboard

Make sure you look at the **Overview**, **Workloads**, and **Pods** views

The screenshot shows the Kubernetes Dashboard interface. The sidebar on the left contains navigation links for Cluster, Namespaces, Nodes, Persistent Volumes, Roles, Storage Classes, Namespace (default), Overview, Workloads (selected), Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, and Stateful Sets. The main content area is titled 'Workloads' and contains three sections:

- Deployments**: A table with columns Name, Labels, Pods, Age, and Images. It shows one deployment named 'nginx-deployment' with 4 pods, labeled 'app: nginx', and an age of 4 minutes.
- Pods**: A table with columns Name, Node, Status, Restarts, and Age. It lists four running pods: 'nginx-deployment-58bd7995b-rq2t8' on worker10, 'nginx-deployment-58bd7995b-6gbsj' on worker12, 'nginx-deployment-58bd7995b-6sntr' on worker10, and 'nginx-deployment-58bd7995b-nshqh' on worker11. All are in 'Running' status with 0 restarts.
- Replica Sets**: A table with columns Name, Labels, Pods, Age, and Images. It shows one replica set named 'nginx-deployment-58bd7995b' with 4 pods, labeled 'app: nginx' and 'pod-template-hash: 146...', and an age of 4 minutes.

## Task 3: Deploy a simple application (Command Line option)

10. To deploy the pod, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSv4/labs/nginx-deployment.yaml**

```
tux@management:~> kubectl apply -f ~/course_files/CAAS101/manifests/labs/nginx-deployment.yaml
2018-03-13 09:27:42.158048 I | proto: duplicate proto type registered: google.protobuf.Any
2018-03-13 09:27:42.158121 I | proto: duplicate proto type registered: google.protobuf.Duration
2018-03-13 09:27:42.158136 I | proto: duplicate proto type registered: google.protobuf.Timestamp
deployment "nginx-deployment" created
tux@management:~>
```

You should see the deployment “nginx-deployment” was created.



11. Enter the following command to view the deployments:

**kubectl get deployments**

You should see that a single instance of the **nginx-deployment** pod is running.

12. Enter the following command to view the deployed pods:

**kubectl get pods**

You should see a single instance of the **nginx-deployment** pod running.

---

### Summary:

In this exercise, you launched a single instance of the Nginx web server as a pod in a deployment on the cluster.

(End of Exercise)

## 3- 2 Delete a Deployment on Kubernetes

---

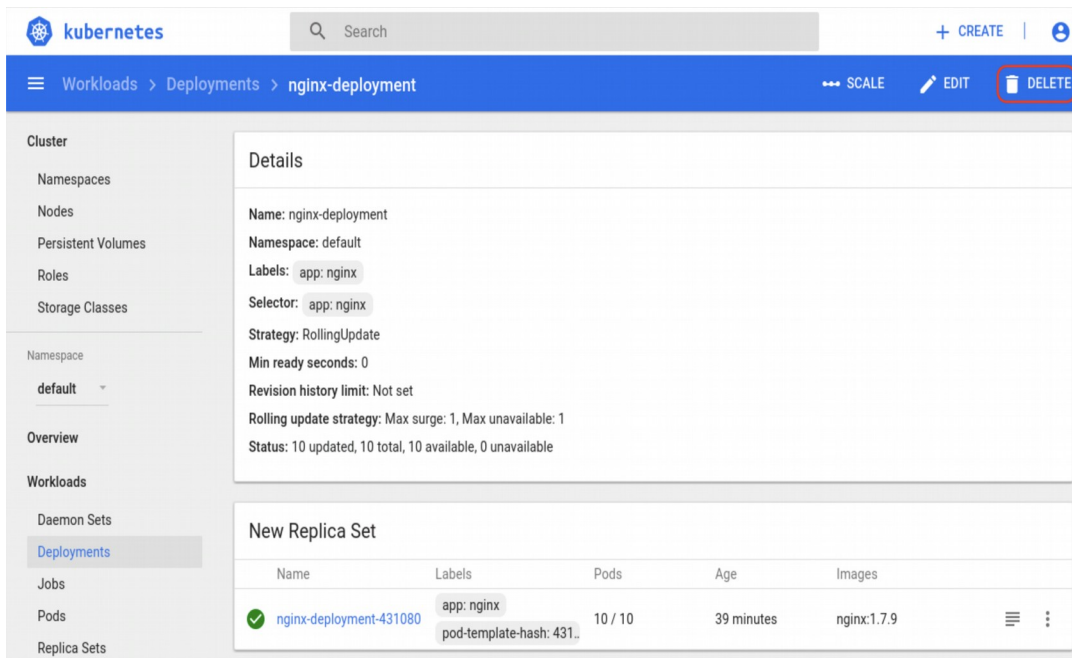
### Description:

In this exercise, you delete the Nginx web server deployment that was previously deployed on the Kubernetes cluster.

---

### Task 1: Delete the Deployment (GUI Option)

1. Select **nginx-deployment** under Deployments
2. Select **Delete**



## Task 2: Delete the Deployment (Command Line option)

- To view the deployments, on the **Workstation**, enter the following command:

**kubectl get deployments**

You should see that one instance of the **nginx-deployment** is running.

- Enter the following command to delete the deployment:

**kubectl delete deployment nginx-deployment**

You should see the **nginx-deployment** was deleted.

- View the deployments again:

**kubectl get deployments**

You should see that the **nginx-deployment** is no longer running.

- View the pods:

**kubectl get pods**

You should see that all of the pods that were part of the **nginx-deployment** are no longer running.

---

## Summary:

In this exercise, you deleted the Nginx web server deployment that was previously deployed on the cluster.

(End of Exercise)

## 3- 3 Scale Out a Deployment

---

### Description:

In this exercise, you scale out a running Deployment.

---

### Task 1: View the Manifest for the Deployment

1. In the text editor of your choice, view the **~/STW-CaaSPv4/labs/nginx-deployment.yaml** file

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 4
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: smt.example.com:5000/sles12sp3_nginx
        ports:
        - containerPort: 80
```

### Task 2: Scale the Deployment (GUI option)

1. From **Kubernetes Dashboard** select **Create** and open **~/STW-CaaSPv4/labs/nginx-scale.yaml**

2. Press the **Deploy** button
3. Select the **nginx-deployment** under **Deployments**
4. click the **Scale** Button and change the number of **Desired pods** to 10

The screenshot shows the Kubernetes dashboard interface. The breadcrumb navigation at the top reads 'Workloads > Deployments > nginx-deployment'. In the top right corner, there are three buttons: 'SCALE' (highlighted with a red box), 'EDIT', and 'DELETE'. The left sidebar contains a navigation menu with categories like Cluster, Namespace, Overview, Workloads, and Config and Storage. Under the 'Workloads' category, 'Deployments' is selected. The main content area displays the details for the 'nginx-deployment' in the 'default' namespace. It shows labels, selector, strategy, and status. Below the details, there is a 'New Replica Set' table with one entry: 'nginx-deployment-431080' with 1 pod. The 'Old Replica Sets' section is empty, and the 'Horizontal Pod Autoscalers' section is also empty.

5. View the deployed pods by selecting **Pods** in the left hand panel

The screenshot shows the Kubernetes dashboard interface with the 'Pods' page selected. The breadcrumb navigation at the top reads 'Workloads > Pods'. The left sidebar shows the 'Pods' option selected under the 'Workloads' category. The main content area displays a table of pods. The table has columns for Name, Node, Status, Restarts, and Age. There are 10 pods listed, all with a status of 'Running' and 0 restarts. The pods are named 'nginx-deployment-431080787-15rg9' through 'nginx-deployment-431080787-if344'.

Name	Node	Status	Restarts	Age
nginx-deployment-431080787-15rg9	b099b425f0194d18902aa1	Running	0	5 minutes
nginx-deployment-431080787-bjd3j	dda6d06e26214c6d9c208	Running	0	5 minutes
nginx-deployment-431080787-ck5g4	b099b425f0194d18902aa1	Running	0	5 minutes
nginx-deployment-431080787-fm6db	e8185390bbc544efba496f	Running	0	5 minutes
nginx-deployment-431080787-fp09k	dda6d06e26214c6d9c208	Running	0	5 minutes
nginx-deployment-431080787-fx8m3	e8185390bbc544efba496f	Running	0	5 minutes
nginx-deployment-431080787-gshzg	b099b425f0194d18902aa1	Running	0	5 minutes
nginx-deployment-431080787-gtmb8	b099b425f0194d18902aa1	Running	0	5 minutes
nginx-deployment-431080787-mdx66	e8185390bbc544efba496f	Running	0	5 minutes
nginx-deployment-431080787-if344	dda6d06e26214c6d9c208	Running	0	12 minutes

6. Scale the Pods back down to 4



### Task 3: Scale the Deployment (Command Line Option)

1. In the text editor of your choice, open the **~/STW-CaaSPv4/labs/nginx-scale.yaml** file

Note the the only difference is the number of replicas

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 10
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: smt.example.com:5000/sles12sp3_nginx
          ports:
            - containerPort: 80
```

2. To scale the deployment, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-scale.yaml**

3. Enter the following command to view the deployments:

**kubectl get deployments**

You should see that one instance of the **nginx-deployment** Deployment is running.

4. Enter the following command to view the deployed pods:

**kubectl get pods**

You should see 10 instances of the **nginx-deployment** pod running.

---

#### Summary:

In this exercise, you created a new manifest for an existing deployment that specified a smaller number of replicas. You then applied the updated manifest to scale in the deployment. Finally you applied the original manifest to scale the deployment back out.

(End of Exercise)

## 4 Work With Kubernetes

---

### Description:

In this section you are introduced to how to work with Kubernetes. You are first introduced to Kubernetes configuration and management utilities and manifests. You then deploy and manage pods on a Kubernetes cluster.

### 4- 1 Update a Deployment

---

### Description:

In this exercise, you update a running pod.

---

### Task 1: Create a New Manifest for the Deployment

1. In the text editor of your choice, open the **~/STW-CaaSPv4/labs/nginx-update.yaml** file
2. Notice the section in **red**

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 4
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: rmt.example.com:5000/nginx:1.9.0
          ports:
            - containerPort: 80
```

## Task 2: Update the Deployment

1. If you don't already have **nginx** deployed Deploy either through **KubeDashboard** using the file **~/STW-CaaSPv4/labs/nginx-pre-update.yaml** or the following command line:

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-pre-update.yaml**

2. To display information on the current nginx deployment enter the following command:

**kubectl describe deployment -l app=nginx**

You should see the description of the nginx deployment displayed.

Notice the image version is: nginx:1.7.9

3. Open another terminal and enter the following command to watch the running pods:

**watch kubectl get pods**

You should see a list of the running pods displayed with the list updating every 2 seconds.

4. To update the deployment, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-update.yaml**

You should see the deployment "nginx-deployment" was configured.

5. Enter the following command to view the deployments:

**kubectl get deployments**

You should see that 4 instances of the **nginx-deployment** pod are DESIRED and, depending on when you ran the command, the values in the CURRENT, UP-TO-DATE and AVAILABLE columns may be more, fewer or the same number as the update happens .

6. In the terminal where you are watching the pods enter **ctrl+c** to stop the watch command
7. Enter the following command to display information about the running deployment of nginx:

**kubectl describe pods -l app=nginx**

Notice the image version is now: nginx:1.9.0

---

### Summary:

In this exercise, you created a new manifest to update the running nginx deployment. You then updated the deployment and verified that it was updated.



## 4- 2 Update a Deployment Via Rolling Updates

---

### Description:

In this exercise, you update a running pod using rolling updates.

---

### Task 1: Create a New Manifest for the Deployment

1. On the **management workstation**, In the text editor of your choice, open the **~/STW-CaaSPv4/labs/nginx-rolling\_update.yaml** file
2. Notice the changes are in **red**

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 4
  revisionHistoryLimit: 5
  minReadySeconds: 20
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 25%
      maxSurge: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: rmt.example.com:/5000/nginx:1.12.0
        ports:
        - containerPort: 80
```

3. Save the file and close the text editor

## Task 2: Update the Deployment

1. To display information on the current nginx deployment enter the following command:

**kubectl describe deployment -l app=nginx**

You should see the description of the nginx deployment displayed.

Notice the image version is: **nginx:1.9.0**

2. Open another terminal and enter the following command to watch the running pods:

**watch kubectl get pods**

You should see a list of the running pods displayed with the list updating every 2 seconds.

3. To update the deployment, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSv4/labs/nginx-rolling\_update.yaml**

You should see the deployment “nginx-deployment” was configured.

In the terminal where you are watching the pods you should see the number of running pods drop to 3 (because of maxUnavailable: 25%) and 3 new pods start deploying. Shortly after the 3 new pods are running you should see the number of running pods scale up to 6 (because of maxSurge: 2) and then back to 4.

4. Enter the following command to view the deployments:

**kubectl get deployments**

You should see that 4 instances of the **nginx-deployment** pod are running.

5. In the terminal where you are watching the pods enter **ctrl+c** to stop the watch command
6. Enter the following command to display information about the running deployment of nginx:

**kubectl describe deployment -l app=nginx**

Notice the image version is now: **nginx:1.12.0**

---

### Summary:

In this exercise, you created a new manifest to update the running nginx deployment using the rolling update type. You then updated the deployment and verified that it was updated.

(End of Exercise)

## 4- 3 Expose a Service Running in a Pod

---

**Description:**

In this exercise, you expose a service running in a pod.

---

**Task 1: Create a Manifest for the Service**

1. On the **Workstation**, in the text editor of your choice, open the file:

**~/STW-CaaSPv4/labs/nginx-service.yaml**

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 30000
  selector:
    app: nginx
```

**Task 2: Define the Service**

1. To define the service in the cluster, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-service.yaml**

You should see the service “nginx-service” was created.

2. Enter the following command to view the services:

**kubectl get services**

You should see that the **nginx-service** service is defined. Notice the 80:30000 under ports showing external port 30000 will be redirected into internal port 80.

**Task 3: Access the Exposed Service**

1. To access the exposed service, open a web browser and point to:

**http://worker.example.com:30000**

You should see the Welcome to nginx web page.



**Note:**

If desired, you can change the URL to point to a specific worker node (i.e. **worker10.example.com**) and see that the web page is still accessible. This demonstrates that the kube-proxy is working on all of the cluster nodes.

---

**Summary:**

In this exercise, you defined a service in the cluster exposing port 80 that allowed access to the nginx pod running on the cluster. You then accessed the nginx pod in a web browser.

(End of Exercise)

## 4- 4 Setup Readiness and Liveness Probes

---

**Description:**

In this exercise, you learn how to setup and configure Readiness and Liveness Probes

---

### Task 1: View manifest for the Deployment

1. In the text editor of your choice, open the **~/STW-CaaSPv4/labs/nginx-deployment-health.yaml** file

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 4
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: smt.example.com:5000/sles12sp3_nginx
          ports:
            - containerPort: 80
          readinessProbe:
            httpGet:
              path: /
              port: 80
            initialDelaySeconds: 5
            timeoutSeconds: 1
            periodSeconds: 15
          livenessProbe:
            httpGet:
              path: /
              port: 80
            initialDelaySeconds: 15
            timeoutSeconds: 1
            periodSeconds: 15
```

2. Notice the ReadinessProbe and LivenessProbe settings

## Task 2: Deploy Manifest

1. Deploy the manifest

**kubectl apply -f ~/STW-CaaSv4/labs/nginx-deployment-health.yaml**

## Task 3: View the Replica sets

2. On the lefthand side of the Kubernetes Dashboard, under Workload, select **Replica Sets**

The screenshot shows the Kubernetes Dashboard interface. The left sidebar has a 'Workloads' section with 'Replica Sets' selected. The main content area shows the details for 'nginx-deployment-6c4cd8dd7'. It includes a 'Pods' table with the following data:

Name	Node	Status	Restarts	Age	CPU (cores)	Memory (bytes)
nginx-deployment-6c4cd8dd7-1	worker12	Running	1	an hour	0	9.359 Mi
nginx-deployment-6c4cd8dd7-2	worker11	Running	0	an hour	0	9.504 Mi
nginx-deployment-6c4cd8dd7-3	worker10	Running	1	an hour	0	9.348 Mi
nginx-deployment-6c4cd8dd7-4	worker10	Running	0	an hour	0	9.355 Mi

Below the Pods table, the 'Services' and 'Horizontal Pod Autoscalers' sections both display 'There is nothing to display here'.

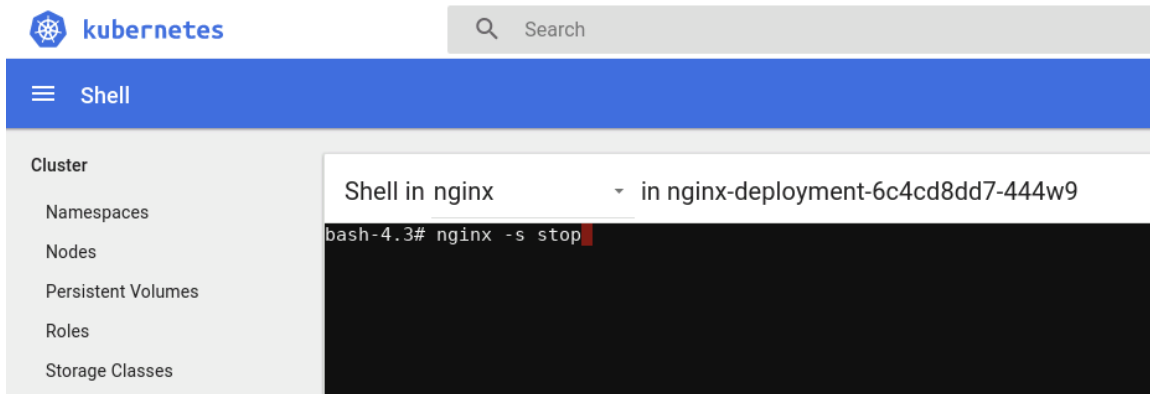
3. Notice the number of Restarts on the container

## Task 4: Kill an nginx server

1. Select one of the pods labeled **nginx-deployment-####** by **clicking** on it
2. **Click** the **EXEC** button to be taken to a shell prompt inside the running pod

The screenshot shows the Kubernetes Dashboard interface. The left sidebar has a 'Workloads' section with 'Pods' selected. The main content area shows the details for 'nginx-deployment-6c4cd8dd7-bmrnp'. At the top, there is an 'EXEC' button. Below it, there are sections for 'CPU usage' and 'Memory usage'.

1. From the Shell enter the following command to kill the nginx process  
**nginx -s stop**



2. Click on Pods to view the current pods running

Overview

Workloads

Cron Jobs

Daemon Sets

Deployments

Jobs

Pods

Replica Sets

Replication Controllers

Pods

Name ↕	Node	Status ↕	Restarts	Age ↕	CPU (cores)	Memory (bytes)
<div><div>✓</div><div>nginx-deployment-6c4cd8dd</div></div>	worker12	Running	2	11 minutes	<div><div></div><div>0</div></div>	<div><div></div><div>9.336 Mi</div></div>
<div><div>✓</div><div>nginx-deployment-6c4cd8dd</div></div>	worker11	Running	0	11 minutes	<div><div></div><div>0</div></div>	<div><div></div><div>9.348 Mi</div></div>
<div><div>✓</div><div>nginx-deployment-6c4cd8dd</div></div>	worker10	Running	0	11 minutes	<div><div></div><div>0</div></div>	<div><div></div><div>9.336 Mi</div></div>
<div><div>✓</div><div>nginx-deployment-6c4cd8dd</div></div>	worker12	Running	0	11 minutes	<div><div></div><div>0</div></div>	<div><div></div><div>9.336 Mi</div></div>

3. Notice the number of restarts has increased. This is because Kubernetes could no longer reach the server on port 80 so it restarted that Nginx Container inside the Pod

## Summary:

In this exercise, you created a new manifest to update the running nginx deployment. You then updated the deployment and verified that it was updated.

(End of Exercise)

## 4- 5 Define Limits for Containers and Pods in Kubernetes

### Description:

In this exercise, you define limits for containers and pods in the Kubernetes cluster.

## Task 1: Create a New Namespace in the Cluster

1. On the **Workstation**, at the command line, enter the following command to create a new namespace in the Kubernetes cluster:

**kubectl create namespace limit-example**

You should see that a new namespace was created.

2. Enter the following command to display the namespaces:

**kubectl get namespaces**

You should see the new namespace listed.

## Task 2: Create a Manifest for the Limits

1. In the text editor open the file:

**~/STW-CaaSv4/labs/limits.yaml**

```
apiVersion: v1
kind: LimitRange
metadata:
  name: mylimits
spec:
  limits:
  - max:
      cpu: "2"
      memory: 1Gi
    min:
      cpu: 200m
      memory: 6Mi
    type: Pod
  - default:
      cpu: 300m
      memory: 200Mi
    defaultRequest:
      cpu: 200m
      memory: 100Mi
    max:
      cpu: "2"
      memory: 1Gi
    min:
      cpu: 100m
      memory: 3Mi
    type: Container
```

## Task 3: Apply the Limits to the Namespace

1. To deploy the pod, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSv4/labs/limits.yaml --namespace=limit-example**

You should see the limitrange “mylimits” was created.

2. Enter the following command to display the limitranges:

**kubectl describe limitranges --namespace=limit-example**

You should see the **mylimits** limitrange listed.



```

Name:      mylimits
Namespace: limit-example
Type:      Resource
-----
Pod        cpu          200m    2        -          -          -
Pod        memory       6Mi     1Gi     -          -          -
Container  cpu          100m    2        200m      300m      -
Container  memory       3Mi     1Gi     100Mi     200Mi     -
tux@management:~/Downloads>

```

**Summary:**

In this exercise, you created a new namespace in the Kubernetes cluster. You then defined limits for containers and pods and applied them to the new namespace.

(End of Exercise)

## 4- 6 Introduction to Helm

**Description:**

In this exercise, you are introduced to the helm utility.

### Task 1: Use Some Basic Helm Commands

1. In a terminal, enter the following command to source the Helm auto-completion functions into your shell environment:

**source <(helm completion bash)**

You shell environment should now have the helm auto-completion functions available.

2. Enter the following command to display a list of available helm charts:

**helm search**

You should see a list of available helm charts listed.

3. Try running the command again but rather than typing in the entire command, only type:

**helm sea[Tab]**

(Where [Tab] is the tab key)

Notice that the “**search**” option is auto-completed.

4. Enter the following command to search for a specific helm chart:

**helm search dokuwiki**

You should see the dokuwiki chart listed.

5. Enter the following command to display the list of currently configured repos:

## helm repo list

You should see a repo for <https://kubernetes-charts.storage.googleapis.com>, [kubernetes-charts.suse.com](https://kubernetes-charts.suse.com) and an addition to the local repo.

---

### Summary:

In this exercise, you initialized helm. You then ran some basic helm commands such as configuring auto-completion, searching for helm charts and displaying repositories.

(End of Exercise)

## 4- 7 Deploy an Application with Helm

---

### Description:

In this exercise, you deploy an application from a helm chart using the helm command.

---

### Task 1: Create Helm Chart Config File

1. On the **Workstation**, in a terminal, enter the following command to search for a dokuwiki helm chart:

#### helm search dokuwiki

You should see a helm chart named stable/dokuwiki listed with its available chart version.

2. Enter the following command to view the default configuration for the dokuwiki chart:

#### helm inspect stable/dokuwiki | less

You should see the configuration displayed in the less pager. Page through the configuration to see what variables are being set.

3. In the text editor of your choice, create/open the **~/STW-CaaSPv4/labs/dokuwiki-config.yaml** file
4. Enter the following in the file:

```
dokuwikiPassword: password123  
serviceType: NodePort  
persistence:  
  enabled: false
```

5. Save the file and close the text editor

## Task 2: Deploy the Helm Chart

1. In a terminal, enter the following command to view the current helm releases:

**helm list**

You should not see the **dokuwiki** application listed.

2. Enter the following command to deploy the chart:

**helm install -f ~/STW-CaaSPv4/labs/dokuwiki-config.yaml --name mywiki stable/dokuwiki**

You should see the chart was deployed.

In the output of the chart deployment, in the **==> v1/Service** section, notice the IP and port(s) the application is listening on. The NodePort(s) that the application is listening on are the number after the colon (:) in the **PORT(S)** column.

Example: **80:32313/TCP,443:31034/TCP**

In this example the NodePorts are **32313** for http and **31034** for https.

Record the http NodePort here:

**DOKUWIKI\_PORT=**\_\_\_\_\_

3. Enter the following command to display the current helm releases:

**helm list**

You should now see the **dokuwiki** chart listed with a name of **mywiki**. Notice the **REVISION** column shows the number **1** as this is the initial deployment of this release.

4. Enter the following command to view the release history for the application:

**helm history mywiki**

You should see that only a single release of **mywiki** exists.

5. Enter the following command to view the status of the **mywiki** release:

**helm status mywiki**

You should see output similar to what was displayed when the chart was first deployed.

6. You can also enter the following **kubectl** commands:

**kubectl get deployments**

**kubectl get pods**

**kubectl get services**

Notice that you see that same info about the deployed deployments/pods/services as if you were to have deployed them from manifests.

## Task 3: Access the Application Deployed by the Helm Chart

1. Open a web browser and point to:

**`http://worker.example.com:DOKUWIKI_PORT`**

You should see the **My Wiki** page displayed.

2. On the top right of the page click: **Login**
3. Enter the following credentials:

**Username:** `user`

**Password:** `password123`

You should be logged in.

After logging in you probably see a number of warnings of available hotfixes and/or updates.

## Task 4: Update the Application Release

1. On the **Workstation**, in a terminal, make a copy of the chart config file you created at the beginning of this exercise:

**`cp ~/dokuwiki-config.yaml ~/dokuwiki-config-update.yaml`**

2. In the text editor of your choice, open the `~/dokuwiki-config-update.yaml` file
3. Edit the file to match the following (changes are in **red**):

**`images: bitnami/dokuwiki:latest`**

**`dokuwikiPassword: password123`**

**`serviceType: NodePort`**

**`persistence:`**

**`enabled: false`**

4. Save the file and close the text editor
5. Enter the following command to upgrade the **mywiki** release:

**`helm upgrade -f ~/dokuwiki-config-update.yaml mywiki stable/dokuwiki`**

You should see the chart was successfully deployed.

6. Enter the following command to view the current helm releases:

**`helm list`**

You should see the **dokuwiki** chart named **mywiki** listed. Notice the **REVISION** column now shows the number **2** as the release has been updated once.

7. Enter the following command to display the release history for the **mywiki** release:

**`helm history mywiki`**

Notice that there are 2 revisions. Also notice the the first revision's **STATUS** is **SUPERSEDED** and the second revision's **STATUS** is **DEPLOYED**.

8. In the web browser, refresh the web page (or log back in if you have logged out)  
You should no longer see the warning messages about available hotfixes and/or

updates.

## Task 5: Delete a Deployed Helm Chart Release

1. In a terminal, enter the following command to delete the **mywiki** release:

**helm delete mywiki**

You should see that the release was deleted.

2. Enter the following command to list the current helm releases:

**helm list**

You should no longer see the **dokuwiki** chart named **mywiki** displayed.

3. Enter the following command to display the status of the **mywiki** release:

**helm status mywiki**

Notice that the **STATUS** is **DELETED** but also that it remembered that it had been deployed as it has a date listed in **LAST DEPLOYED**.

---

### Summary:

In this exercise, you first deployed a helm chart. You then accessed the application that was deployed, Next you updated the release, verifying it was updated. Finally you deleted the release.

(End of Exercise)

## 5 Deploying Advanced Workload

---

### Description:

How to deploy a workload from an existing Docker Image

### 5- 1 Run Heimdall as a standalone container

---

### Description:

---

In this exercise, you configure and run Heimdall as standalone container

---

### Task 1: Create a place for Heimdall to store it's files

1. On the **Workstation**, enter the following commands in a terminal session to create the `~/docker/heimdall` folder:

```
sudo mkdir -p /docker/heimdall  
sudo chmod 777 /docker -R  
cd /docker/heimdall
```

The directory and file should now exist.

### Task 2: Run Heimdall as a container

1. On the **Workstation**, enter the following commands in a terminal session to create

```
sudo docker run -p 80:80 -v /docker/heimdall:/config -e  
TZ=Europe/London smt.example.com:5000/heimdall
```

This tells docker to run the heimdall container with the following options:

**Map port 80 in the container to 80 on the host**

**mount `~/docker/heimdall` in the container in `/config`**

**Set the timezone within the container**

2. Notice that the first time it runs it created a default installation



Brought to you by linuxserver.io  
 We gratefully accept donations at:  
<https://www.linuxserver.io/donate/>

GID/UID

User uid: 911  
 User gid: 911

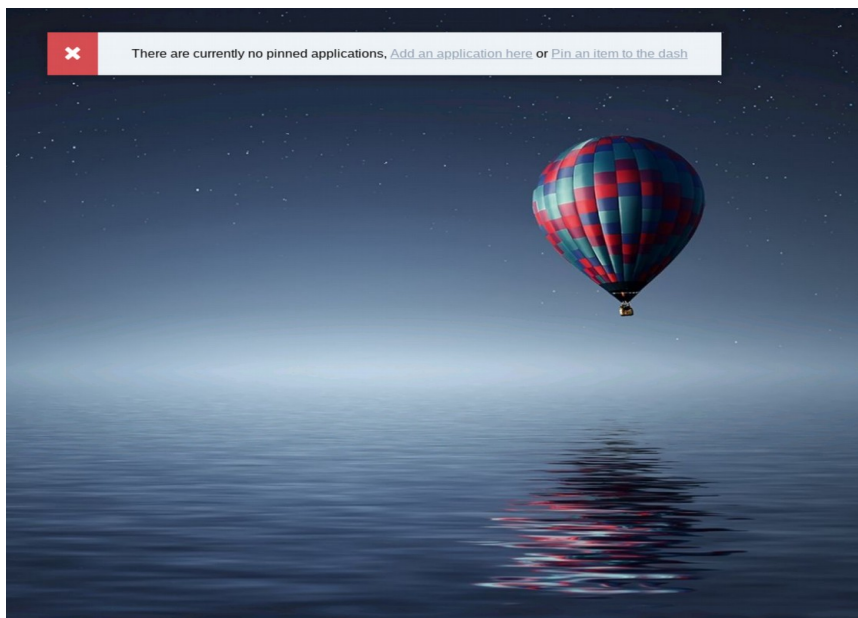
```
[cont-init.d] 10-adduser: exited 0.
[cont-init.d] 20-config: executing...
[cont-init.d] 20-config: exited 0.
[cont-init.d] 30-keygen: executing...
generating self-signed keys in /config/keys, you can rep
Generating a RSA private key
.....+++++
.....
writing new private key to '/config/keys/cert.key'
-----
[cont-init.d] 30-keygen: exited 0.
[cont-init.d] 50-config: executing...
New container detected, installing Heimdall
Creating app key. This may take a while on slower system
Application key set successfully.
Setting permissions
[cont-init.d] 50-config: exited 0.
[cont-init.d] 99-custom-scripts: executing...
[custom-init] no custom scripts found exiting...
[cont-init.d] 99-custom-scripts: exited 0.
[cont-init.d] done.
[services.d] starting services
[services.d] done.
```

### Task 3: On the Workstation launch the Chrome Browser

1. On the **Workstation** launch the **Chrome Browser**

<http://127.0.0.1>

2. You should see a default screen with no Apps Defined



## Task 4: Create an App in Helm

1. Click on 'Add an application here'

Application Name: **SUSE**

URL: <http://www.suse.com>

Select: **Pinned**

A screenshot of the "Add application" form in the SUSE CaaS Platform. The form is titled "Add application" and has a "PINNED" toggle switch, a "SAVE" button, and a "CANCEL" button. The form fields are: "Application name \*" with the value "SUSE", "Application Type \*" with a dropdown menu showing "None", "Colour \*" with a "Hex colour" input field, "URL" with the value "http://www.suse.com", and "Tags (Optional)" with an empty input field. There is also an "Upload a file" button next to a circular icon. Below the form is a "Preview" section showing a card with the SUSE logo and the name "SUSE". At the bottom right of the form are "SAVE" and "CANCEL" buttons.

2. Click Save to Save the App



## Task 5: Stop and restart container

1. **On the Workstation**, go back to the terminal session running the container and press the following Keys to stop the container instance

**<ctrl>C**

This will terminate the container instance

2. Restart the Docker Container

**sudo docker run -p 80:80 -v ~/docker/heimdall:/config -e TZ=Europe/London smt.example.com:5000/heimdall**

Notice How it start also immediately because it's using the configuration from the last time we launched the Application

## Task 6: Verify with Browser

1. **On the Workstation launch the Chrome Browser**

<http://12.0.0.1>

2. You should see the App you previously defined

## Task 7: Stop Container Instance

1. **On the Workstation**, go back to the terminal session running the container and press the following Keys to stop the container instance

**<ctrl>C**

This will terminate the container instance

## Task 8: Try and run multiple versions of Heimdall

1. **On the Workstation**, go back to the terminal session running the container and press the following Keys to stop the container instance

**sudo docker run -d -p 80:80 -v ~/docker/heimdall:/config -e TZ=Europe/London smt.example.com:5000/heimdall**

**\* notice we added the -d command to tell docker to run as a Daemon rather than in the foreground**

1. Try and run a second instance of the Heimdall

```
sudo docker run -d -p 80:80 -v ~/docker/heimdall:/config -e  
TZ=Europe/London smt.example.com:5000/heimdall
```

**Notice is fails because the port is already in use**

---

### Summary:

In this exercise, you launched a container from the command line to see how it behaved

(End of Exercise)

## 5- 2 Launch Heimdall

---

### Description:

In this exercise, we will Heimdall without dedicated storage

---

### Task 1: Create the Manifest for the Persistent Volume

1. On the **Workstation**, in the text editor of your choice, open the file:

```
~/STW-CaaSPv4/labs/heimdall/app-health/heimdall-deployment_ns.yaml
```

Enter the following in the file:

```

apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: heimdall-shared-deployment
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: heimdall-shared
    spec:
      containers:
        - name: heimdall-pod
          image: linuxserver/heimdall
          ports:
            - containerPort: 80

```

2. On the **Workstation**, in the text editor of your choice, open the file:

**~/STW-CaaSPv4/labs/heimdall/app-health/heimdall-service.yaml**

Enter the following in the file:

```

apiVersion: v1
kind: Service
metadata:
  name: heimdall-shared-service
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 33000
  selector:
    app: heimdall-shared

```

## Task 2: Deploy Heimdall

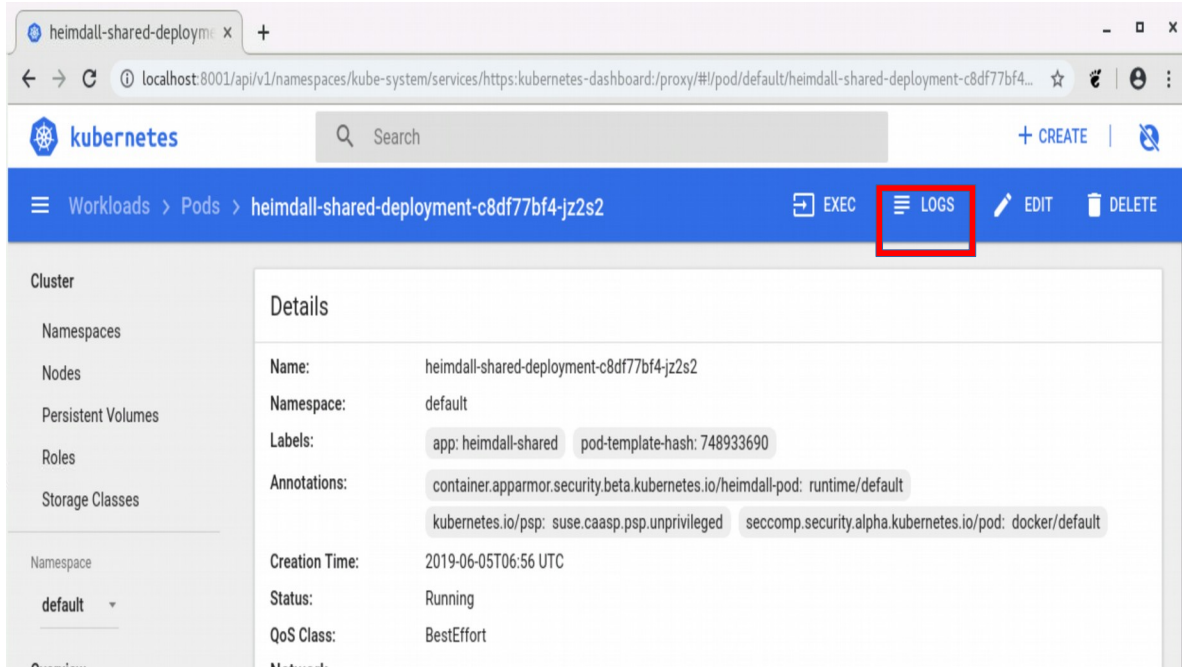
1. From a Terminal prompt on the Management workstation

**kubectl apply -f ~/STW-CaaSPv4/labs/heimdall/app-health/heimdall-deployment\_ns.yaml**

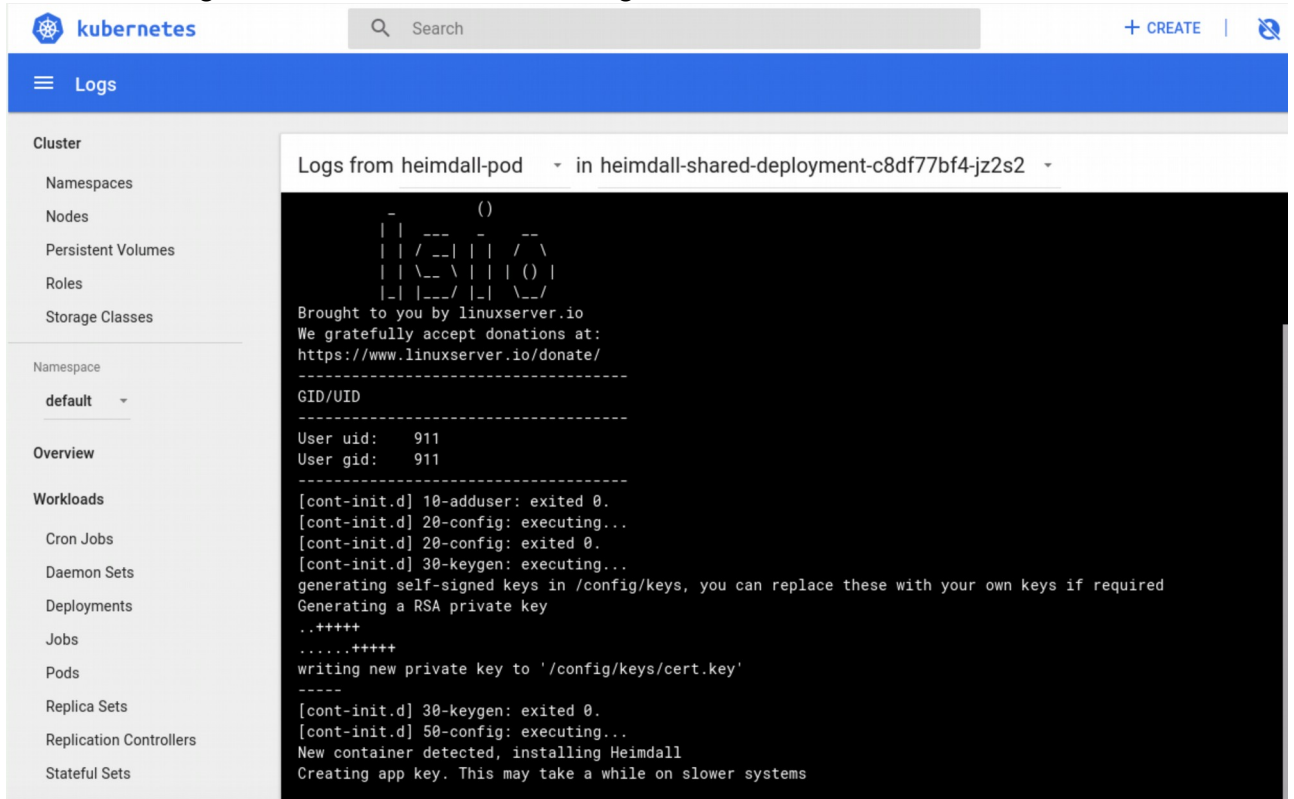
**kubectl apply -f ~/STW-CaaSPv4/labs/heimdall/app-health/heimdall-service.yaml**

## Task 3: View the Deployment

1. From Kubernetes Dashboard select Pods and then click on the Pod that was just deployed
2. Click on the Log button so we can see what is happening on that Pod



- ### 3. Look at the logs and notice how it is creating a new installation



4. Let that deployment continue to run until it is completely finished (Refresh the Browser to check Status)

## Task 4: View Heimdall in a Browser

1. On the **Workstation** open the **Chrome browser** and go to **worker.example.com:33000**
2. Notice how ever time it is deployed it has to rebuild the installation
3. Stop the deployment by either killing it on the Kubernetes Dashboard from from the command line by typing in the following command

```
kubecti delete -f ~/STW-CaaSPv4/labs/heimdall/app-health/heimdall-deployment_ns.yaml
```

## Task 4: Re-deploy Heimdall

1. Repeat all of the steps in **Task 2 and Task 3**
2. Notice how ever time it is deployed it has to rebuild the installation

---

### Summary:

In this exercise, you created manifests for a Heimdall. We them deployed it and watched it build the default configuration. We learned that every time we deployed the app it had to create the default installation.

(End of Exercise)

## 5- 3 Configure Heimdall to use NFS Persistent Storage

---

### Description:

In this exercise, you configure a persistent volume on an NFS server. You then launch a pod that use the persistent volume

---

## Task 1: Create the Manifest for the Persistent Volume

1. On the **Workstation**, in the text editor of your choice, open the file:  
**~/course\_files/CAAS101/course\_files/manifests/labs/heimdall/app-shared/heimdall-pv-shared.yaml**

Enter the following in the file:

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: heimdall-shared
spec:
  capacity:
    storage: 200Mi
  accessModes:
    - ReadWriteMany
  nfs:
    server: 192.168.110.2
    path: "/export/heimdall-shared"
```

## Task 2: Create the Manifest for the Persistent Volume Claim

1. In the text editor of your choice, open the file:

**~/course\_files/CAAS101/course\_files/manifests/labs/heimdall/app-shared/heimdall-pvc-shared.yaml**

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: heimdall-shared-claim
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 200Mi
```

## Task 3: Create the Manifest for the Service so we can access the deployment

1. In the text editor of your choice, open the file:
2. **~/course\_files/CAAS101/course\_files/manifests/labs/heimdall/app-shared/heimdall-service.yaml**

```
apiVersion: v1
kind: Service
metadata:
  name: heimdall-shared-service
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 30000
  selector:
    app: heimdall-shared
```

#### Task 4: Create the Manifests for the Web Frontend

1. In the text editor of your choice, create/open the file:
2. **~/course\_files/CAAS101/course\_files/manifests/labs/heimdall/app-shared/heimdall-deployment.yaml**



```

apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: heimdall-shared-deployment
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: heimdall-shared
    spec:
      containers:
        - name: heimdall-pod
          image: linuxserver/heimdall
          volumeMounts:
            - name: heimdall-pvc-shared
              mountPath: "/config"
      volumes:
        - name: heimdall-pvc-shared
          persistentVolumeClaim:
            claimName: heimdall-shared-claim

```

## Task 5: Deploy the Objects

1. To deploy the volumes/pods/service, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/heimdall/app-shared**

You should see the following were created (not necessarily in this order):

```

deployment.extensions "heimdall-shared-deployment" created
persistentvolume "heimdall-shared" created
persistentvolumeclaim "heimdall-shared-claim" created
service "heimdall-shared-service" created

```

2. Enter the following command to view the deployments:

**kubectl get deployments**

You should see the **heimdall-shared-deployment** deployments listed.

3. Enter the following command to view the persistent volumes:

```
kubectl get pv
```

You should see the persistent volume **heimdall-shared** listed.

4. Enter the following command to view the persistent volumes:

```
kubectl get pvc
```

You should see the persistent volume claim **heimdall-shared-claim** listed.

## Task 6: Test the Persistent Data

1. On the **management workstation**, open a web browser and point to:

```
http://worker10.example.com:30000
```

Notice we have a Heimdall session pre-populated with with 2 Applications

## Task 7: Remove the Objects from the Cluster

1. In the first terminal, enter the following commands to delete all of the objects:

```
kubectl delete -f ~/STW-CaaSv4/labs/heimdall/app-shared
```

You should see that the objects were deleted.

---

### Summary:

In this exercise, you created manifests for a persistent volume, a persistent volume claim, a pod to that attached to the volume and writes data to an index.html file in the volume, and a web server that attaches to the volume and displays the index.html file. You then verified that the index.html file was being updated by looking at the file both on the NFS volume and the web server.

(End of Exercise)