

# **Technical Workshop - SUSE CaaS Platform 4 Beta**

## **-Workbook-**

**Course ID: Technical Workshop – SUSE CaaSP**

**Version: RC1**

**Date: 2019-08-21**



## **Proprietary Statement**

Copyright © 2019 SUSE LLC. All rights reserved.

SUSE LLC, has intellectual property rights relating to technology embodied in the product that is described in this document.

No part of this publication may be reproduced, photocopied, stored on a retrieval system, or transmitted without the express written consent of the publisher.

SUSE  
Maxfeldstrasse 5  
90409 Nuremberg  
Germany  
[www.suse.com](http://www.suse.com)

(C) 2019 SUSE LLC. All Rights Reserved. SUSE and the SUSE logo are registered trademarks of SUSE LLC in the United States and other countries. All third-party trademarks are the property of their respective owners.

If you know of illegal copying of software, contact your local Software Antipiracy Hotline.

## **Disclaimer**

SUSE LLC, makes no representations or warranties with respect to the contents or use of this documentation, and specifically disclaims any express or implied warranties of merchantability or fitness for any particular purpose.

Further, SUSE LLC, reserves the right to revise this publication and to make changes to its content, at any time, without obligation to notify any person or entity of such revisions or changes. Further, SUSE LLC, makes no representations or warranties with respect to any software, and specifically disclaims any express or implied warranties of merchantability or fitness for any particular purpose. Further, SUSE LLC, reserves the right to make changes to any and all parts of SUSE software, at any time, without any obligation to notify any person or entity of such changes.

Any products or technical information provided under this Agreement may be subject to U.S. export controls and the trade laws of other countries. You agree to comply with all export control regulations and to obtain any required licenses or classification to export, re-export or import deliverables. You agree not to export or re-export to entities on the current U.S. export exclusion lists or to any embargoed or terrorist countries as specified in the U.S. export laws. You agree to not use deliverables for prohibited nuclear, missile, or chemical biological weaponry end uses. SUSE assumes no responsibility for your failure to obtain any necessary export approvals.

This SUSE Training Manual is published solely to instruct students in the use of SUSE networking software. Although third-party application software packages may be used in SUSE training courses, this is for demonstration purposes only and shall not constitute an endorsement of any of these software applications.

Further, SUSE LLC does not represent itself as having any particular expertise in these application software packages and any use by students of the same shall be done at the student's own risk.

# Table of Contents

Documentation Conventions:	6
<b>Section 1 : Setup Workstation to Initialize Cluster.....</b>	<b>7</b>
Exercise 1 : Start the Lab VMs.....	8
Task 1: Start the Lab VMs.....	8
Exercise 2 : Prepare the Management Workstation.....	10
Task 1: Install needed software on Management Workstation.....	10
Task 2: Test to ensure SSH keys are properly set.....	11
Exercise 3 : Setup SSH Keys (optional).....	12
Task 1: Open a terminal.....	12
Task 2: Create SSH key (This has been done – information only).....	12
Task 3: Add the SSH key to the agent.....	12
Task 4: Test to ensure SSH keys are properly set.....	12
Exercise 4 : Install CaaS Platform Nodes.....	14
Task 1: Initialize the and bootstrap the cluster.....	14
Task 2: Deploy the Worker nodes.....	16
<b>Section 2 : Connecting to the Cluster.....</b>	<b>17</b>
Exercise 1 : Use the kubectl Command to Display Info About the Cluster.....	18
Task 1: Configure workstation for kubectl command.....	18
Task 2: Use the kubectl Command.....	18
Exercise 2 : Deploy the Kubernetes Dashboard.....	21
Task 1: Deploy the Kubernetes Dashboard.....	21
Task 2: Access the Kubernetes Dashboard.....	22
Exercise 3 : Install and Configure Helm.....	24
Task 1: Install Helm on Workstation.....	24
Task 2: Create Service account and Cluster Role for Helm.....	24
Task 3: Initialize Helm.....	24
Task 4: Add SUSE Charts to Helm.....	25
Task 5: Install Centralize Logging.....	26
<b>Section 3 : Deploying a Workload.....</b>	<b>27</b>
Exercise 1 : Deploy a Simple Pod on Kubernetes.....	28
Task 1: View a Manifest for the Deployment.....	28
Task 2: Deploy a simple application (GUI option).....	29
Task 3: Deploy a simple application (Command Line option).....	31
Exercise 2 : Delete a Deployment on Kubernetes.....	33
Task 1: Delete the Deployment (GUI Option).....	33
Task 2: Delete the Deployment (Command Line option).....	33
Exercise 3 : Scale Out a Deployment.....	35
Task 1: View the Manifest for the Deployment.....	35
Task 2: Scale the Deployment (GUI option).....	35
Task 3: Scale the Deployment (Command Line Option).....	37
<b>Section 4 : Working With Kubernetes.....</b>	<b>38</b>

Exercise 1 : Update a Deployment.....	38
Task 1: Create a New Manifest for the Deployment.....	38
Task 2: Update the Deployment.....	39
Exercise 2 : Update a Deployment Via Rolling Updates.....	41
Task 1: Create a New Manifest for the Deployment.....	41
Task 2: Update the Deployment.....	42
Exercise 3 : Expose a Service Running in a Pod.....	42
Task 1: Create a Manifest for the Service.....	43
Task 2: Define the Service.....	43
Task 3: Access the Exposed Service.....	43
Exercise 4 : Setup Readiness and Liveness Probes.....	44
Task 1: View manifest for the Deployment.....	44
Task 2: Deploy Manifest.....	46
Task 3: View the Replica sets.....	46
Task 4: Kill an nginx server.....	46
Exercise 5 : Define Limits for Containers and Pods in Kubernetes.....	47
Task 1: Create a New Namespace in the Cluster.....	48
Task 2: Create a Manifest for the Limits.....	48
Task 3: Apply the Limits to the Namespace.....	48
Exercise 6 : Introduction to Helm.....	49
Task 1: Use Some Basic Helm Commands.....	49
Exercise 7 : Deploy an Application with Helm.....	50
Task 1: Create Helm Chart Config File.....	50
Task 2: Deploy the Helm Chart.....	51
Task 3: Access the Application Deployed by the Helm Chart.....	51
Task 4: Update the Application Release.....	52
Task 5: Delete a Deployed Helm Chart Release.....	53
<b>Section 5 : Storage.....</b>	<b>53</b>
Exercise 1 : Configure NFS Persistent Storage.....	54
Task 1: Create the Persistent Volume on the NFS Server.....	54
Task 2: Create the Manifest for the Persistent Volume.....	54
Task 3: Create the Manifest for the Persistent Volume Claim.....	54
Task 4: Create the Manifest for the Busybox Instance.....	55
Task 5: Create the Manifests for the Web Frontend.....	55
Task 6: Deploy the Objects.....	56
Task 7: Test the Persistent Data.....	57
Task 8: Remove the Objects from the Cluster.....	58
Exercise 2 : Configure Persistent Storage with a NFS StorageClass.....	59
Task 1: Create the Directory for the Persistent Volumes on the NFS Server.....	59
Task 2: Define a Service Account, Cluster Role and Cluster Role Binding.....	59
Task 3: Define a Storage Class.....	61
Task 4: Create a Persistent Volume Claim and Pod that Uses It.....	62
Task 5: Create a Persistent Volume Claim and Pod that Uses It.....	64
Task 6: Remove the Objects from the Cluster.....	65
<b>Section 6 : Deploying Advanced Workload.....</b>	<b>66</b>
Exercise 1 : Run Heimdall as a standalone container.....	66

<b>SUSE U Advanced - SUSE CaaS Platform</b>	
Task 1: Create a place for Heimdall to store it's files.....	66
Task 2: Run Heimdall as a container.....	66
Task 3:On the Workstation launch the Chrome Browser.....	67
Task 4: Create an App in Helm.....	68
Task 5: Stop and restart container.....	69
Task 6: Verify with Browser.....	69
Task 7: Stop Container Instance.....	69
Task 8: Try and run multiple versions of Heimdall.....	69
<b>Exercise 2 : Launch Heimdall.....</b>	<b>70</b>
Task 1: Create the Manifest for the Persistent Volume.....	70
Task 2: Deploy Heimdall.....	71
Task 3: View the Deployment.....	72
Task 4: View Heimdall in a Browser.....	74
Task 4: Re-deploy Heimdall.....	74
<b>Exercise 3 : Configure Heimdall to use NFS Persistent Storage.....</b>	<b>74</b>
Task 1: Create the Manifest for the Persistent Volume.....	74
Task 2: Create the Manifest for the Persistent Volume Claim.....	75
Task 3: Create the Manifest for the Service so we can access the deployment.....	75
Task 4: Create the Manifests for the Web Frontend.....	76
Task 5: Deploy the Objects.....	77
Task 6: Test the Persistent Data.....	78
Task 7: Remove the Objects from the Cluster.....	78

## Documentation Conventions:

---

The following typographical conventions are used in this manual:

<b>Bold</b>	Represents things you should pay attention to or buttons you click, text or options that you should click/select/type in a GUI.
<b>Bold Gray</b>	Represents the name of a Task or in the context of what is seen on the screen, the screen name, a tab name, column name, field name, etc.
<b>Bold Red</b>	Represents warnings or very important information.
<b>Option &gt; Option &gt; Option</b>	Represents a chain of items selected from a menu.
<b><i>BOLD_UPPERCASE_ITALIC</i></b>	Represents an “exercise variable” that you replace with another value.
<b>bold monospace</b>	Represents text displayed in a terminal or entered in a file.
<b>bold monospace blue</b>	Represents commands entered at the command line.
<b>bold monospace green</b>	Represents a file name.

# 1 Setup Workstation to Initialize Cluster

---

## Description:

You will setup a workstation to initialize the cluster

## 1-1 Start the Lab VMs

### Description:

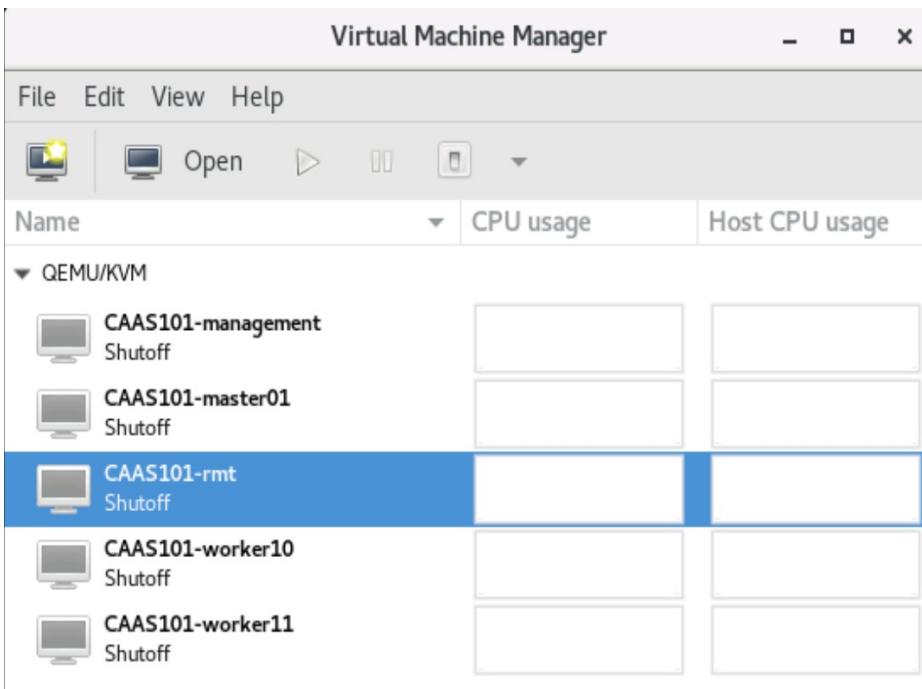
In this exercise, you use the Virt-Manager utility to start the lab VM(s) in the proper order.

### Task 1: Start the Lab VMs

1. On the VM host, launch the Virt-Manager utility  
(4<sup>th</sup> icon from the left on the dock at the bottom of the screen)  
You should see the course VMs listed.



2. Right-click on the CAAS101-rmt VM and select Run



3. Double-click on the CAAS101-rmt VM to view its console

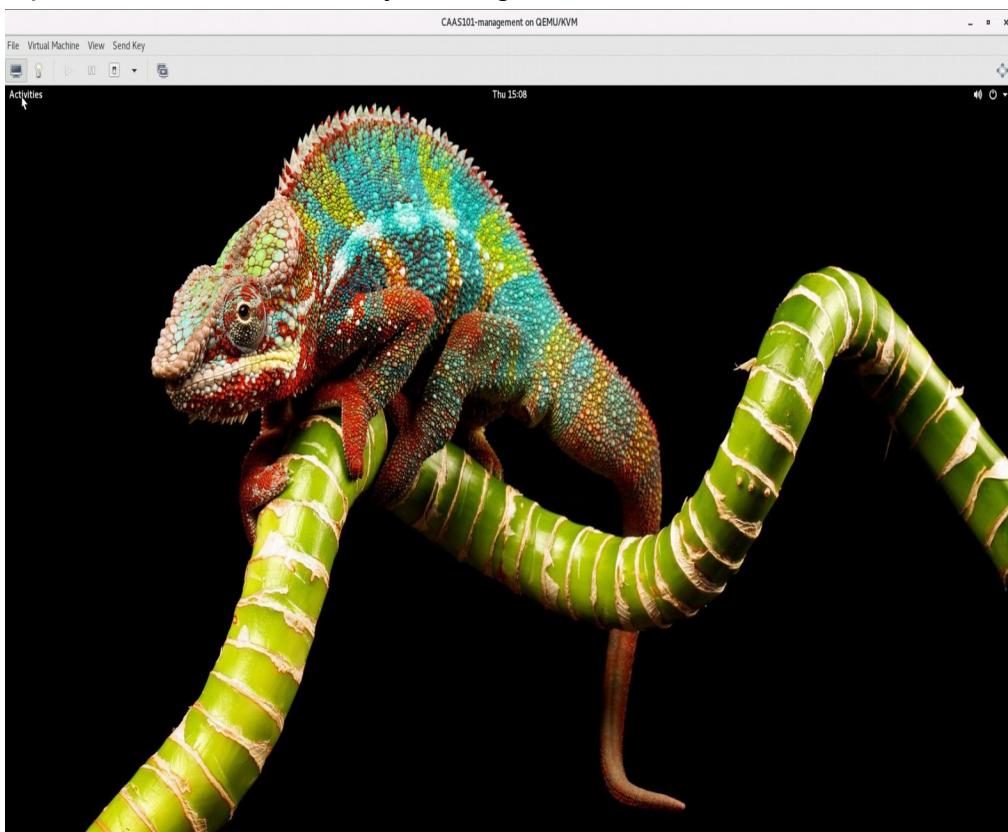
When the CAAS101-rmt VM has finished booting completely, close its console window.  
(You can use the CPU utilization graph in Virt-Manager to see when the node is finished booting and starting services).

4. Power on the reset of the Virtual Machines **CAAS101-management**, **CAAS101-master**, **CAAS101-worker10** and **CAAS101-worker10**

\*note you can power all of the on at the same time

When the **Workstation** has finished booting you will use it as your management workstation for all tasks including managing VMs with the Virt-Manager utility.

5. Make sure you double click on the **CAAS101-management** so you view it full screen
6. Open a Terminal session by clicking on the 3<sup>rd</sup> icon down



---

### **Summary:**

In this exercise, you started the required lab VM(s) in the proper order.

(End of Exercise)

## 1- 2 Prepare the Management Workstation

---

### Description:

In this exercise, you prepare the Management Workstation for the CaaS Platform installation.

### Dependencies:

The RMT server must be configured (Optional Docker Registry)

DHCP, PXE and DNS must be configured.

---

### Task 1: Install needed software on Management Workstation

1. On the **Workstation**, open a terminal
2. Enter the following command to install skuba and the kubernetes client

**sudo zypper in -t pattern SUSE-CaaSP-Management**

It will tell you the amount of additional drive space needed and it will prompt you to press **y** to continue

3. Press the '**Y**' key to continue

when it finished you should see a screen similar to the image below

```
tux@workstation:~ x
File Edit View Search Terminal Help
Retrieving: kubernetes-common-1.15.0-2.4.1.x86_64.rpm ..... [done]
Retrieving package skuba-0.8.1-2.10.2.x86_64
(2/5), 14.4 MiB ( 45.5 MiB unpacked)
Retrieving: skuba-0.8.1-2.10.2.x86_64.rpm ..... [done (10.7 MiB/s)]
Retrieving package skuba-update-0.8.1-2.10.2.noarch
(3/5), 54.4 KiB ( 33.5 KiB unpacked)
Retrieving: skuba-update-0.8.1-2.10.2.noarch.rpm ..... [done]
Retrieving package kubernetes-client-1.15.0-2.4.1.x86_64
(4/5), 1.2 MiB (743.3 KiB unpacked)
Retrieving: kubernetes-client-1.15.0-2.4.1.x86_64.rpm ..... [done]
Retrieving package kubectl-caasp-0.8.1-2.10.2.x86_64
(5/5), 14.4 MiB ( 45.4 MiB unpacked)
Retrieving: kubectl-caasp-0.8.1-2.10.2.x86_64.rpm ..... [done]
Checking for file conflicts: ..... [done]
(1/5) Installing: kubernetes-common-1.15.0-2.4.1.x86_64 ..... [done]
(2/5) Installing: skuba-0.8.1-2.10.2.x86_64 ..... [done]
(3/5) Installing: skuba-update-0.8.1-2.10.2.noarch ..... [done]
Additional rpm output:
Updating /etc/sysconfig/skuba-update ...
(4/5) Installing: kubernetes-client-1.15.0-2.4.1.x86_64 ..... [done]
(5/5) Installing: kubectl-caasp-0.8.1-2.10.2.x86_64 ..... [done]
tux@workstation:~> |
```

## Task 2: Test to ensure SSH keys are properly set

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to test your **ssh keys** to master01  
**ssh root@master01**
3. You should see a red prompt that says master01

```
tux@workstation:~  
File Edit View Search Terminal Help  
tux@workstation:~> ssh root@master01  
Last login: Sat Aug  3 08:54:16 2019 from 192.168.110.108  
master01:~ #
```

4. Notice that you did not need to type in a password or accept a key?  
This is because of the pre-work done on the workstation
5. Exit your ssh session from the terminal prompt  
**exit**
6. Repeat the step 2 and 3 for **worker10** and **worker11**

---

### Summary:

In this exercise, you installed Management workstation utilities. You also tested connectivity to all of the nodes.

(End of Exercise)

## 1-3 Setup SSH Keys (optional)

---

### Description:

In this exercise, you will setup all nodes to ssh using key authentication

This is an **OPTIONAL** exercise to be performed only if you do not have SSH setup on your workstation intended to manage the cluster installation

---

### Task 1: Open a terminal

1. On the **Workstation**, open a terminal

### Task 2: Create SSH key (This has been done – information only)

1. On the **Workstation**, in the terminal opened previously

Enter the following command to create a ssh key

**ssh-keygen**

When it asks you questions just take the defaults

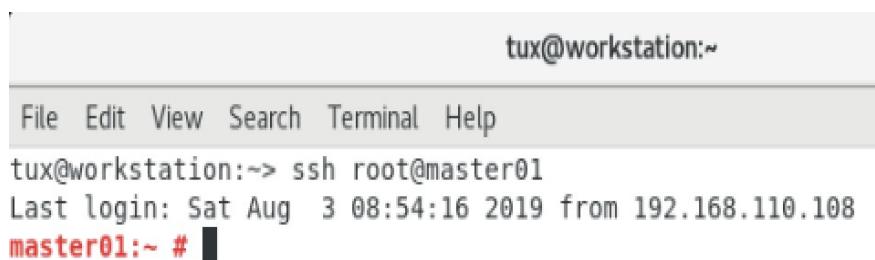
### Task 3: Add the SSH key to the agent

1. On the **Workstation**, in the terminal opened previously
2. Enter the following command in the terminal to add your **ssh key** to ssh-agent

**ssh-add**

### Task 4: Test to ensure SSH keys are properly set

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to test your **ssh keys** to master01  
**ssh root@master01**
3. You should see a red prompt: **master01:~ #**



tux@workstation:~

---

File Edit View Search Terminal Help

tux@workstation:~> ssh root@master01

Last login: Sat Aug 3 08:54:16 2019 from 192.168.110.108

**master01:~ #**

4. Notice that you did not need to type in a password or accept a key?

This is because of the pre-work done on the workstation

5. Exit your ssh session from the terminal prompt

**exit**

6. Repeat step 2 and 3 for **worker10** and **worker11**

---

**Summary:**

In this exercise, you set up and verified the SSH key to ssh from the Management Workstation to any of the nodes using key authentication.

(End of Exercise)

## 1 - 4 Install CaaS Platform Nodes

---

### Description:

In this exercise, you install SUSE CaaS Platform all nodes.

### Dependencies:

The RMT server must be configured (Optional Docker Registry)

DHCP, PXE and DNS must be configured.

Master, and 2 worker nodes with SLES 15 SP1 installed and CaaS Platform Repos added

---

### Task 1: Initialize the and bootstrap the cluster

1. On the **Workstation**, open a terminal
2. Enter the following command create the basic config to our **kubernetes cluster**

**skuba cluster init --control-plane master.example.com my-cluster**

```
tux@workstation:~> skuba cluster init --control-plane master.example.com my-cluster
** This is a BETA release and NOT intended for production usage. **
[init] configuration files written to /home/tux/my-cluster
tux@workstation:~> █
```

3. Enter the following command to change into the newly created folder

**cd my-cluster**

4. Look at the cluster config filename **my-cluster**

**less kubeadm-init.conf**

```
tux@workstation:~/my-cluster> cat kubeadm-init.conf
apiVersion: kubeadm.k8s.io/v1beta1
kind: InitConfiguration
bootstrapTokens: []
localAPIEndpoint:
  advertiseAddress: ""
---
apiVersion: kubeadm.k8s.io/v1beta1
kind: ClusterConfiguration
apiServer:
  certSANs:
    - master.example.com
extraArgs:
  oidc-issuer-url: https://master.example.com:32000
  oidc-client-id: oidc
  oidc-ca-file: /etc/kubernetes/pki/ca.crt
  oidc-username-claim: email
  oidc-groups-claim: groups
clusterName: my-cluster
controlPlaneEndpoint: master.example.com:6443
dns:
  imageRepository: registry.suse.com/caasp/v4
  imageTag: 1.3.1
  type: CoreDNS
etcd:
  local:
    imageRepository: registry.suse.com/caasp/v4
    imageTag: 3.3.11
imageRepository: registry.suse.com/caasp/v4
kubernetesVersion: 1.15.0
networking:
  podSubnet: 10.244.0.0/16
  serviceSubnet: 10.96.0.0/12
useHyperKubeImage: true
tux@workstation:~/my-cluster> █
```

5. Install, configure, and bootstrap the first node as the master

From the **~/my\_cluster** directory:

**skuba node bootstrap --target master.example.com master01**

```
tux@workstation:~/my-cluster> skuba node bootstrap --target master.example.com master01
** This is a BETA release and NOT intended for production usage. **
W0808 17:13:25.590213    2588 ssh.go:306]
The authenticity of host '192.168.110.7:22' can't be established.
ECDSA key fingerprint is 54:3e:9b:2e:21:fc:5c:70:ff:c4:99:7c:fa:b4:08:d0.
I0808 17:13:25.590316    2588 ssh.go:307] accepting SSH key for "master.example.com:22"
I0808 17:13:25.590327    2588 ssh.go:308] adding fingerprint for "master.example.com:22" to "known_hosts"
[bootstrap] updating init configuration with target information
[bootstrap] writing init configuration for node
[bootstrap] applying init configuration to node
[bootstrap] downloading secrets from bootstrapped node "master.example.com"
[bootstrap] successfully bootstrapped node "master.example.com" with Kubernetes: "1.15.0"
tux@workstation:~/my-cluster> █
```

## Task 2: Deploy the Worker nodes

1. Add a **worker** node to the cluster

From the **~/my\_cluster** directory:

**skuba node join --role worker --target worker10.example.com worker10**

```
tux@workstation:~/my-cluster> skuba node join --role worker --target worker10.example.com worker10
** This is a BETA release and NOT intended for production usage. **
W0808 17:23:50.580867    2632 ssh.go:306]
The authenticity of host '192.168.110.10:22' can't be established.
ECDSA key fingerprint is 19:7c:db:38:d2:c5:34:f6:cb:90:21:4e:1a:b5:5e:d7.
I0808 17:23:50.580863    2632 ssh.go:307] accepting SSH key for "worker10.example.com:22"
I0808 17:23:50.580916    2632 ssh.go:308] adding fingerprint for "worker10.example.com:22" to "known_hosts"
[join] applying states to new node
[join] node successfully joined the cluster
tux@workstation:~/my-cluster>
```

2. Add a second **worker** node to the cluster

**skuba node join --role worker --target worker11.example.com worker11**

You should see a message very similar to the message you received in the previous step.

3. Verify cluster status

**skuba cluster status**

```
tux@workstation:~/my-cluster> skuba cluster status
** This is a BETA release and NOT intended for production usage. **
NAME      OS-IMAGE          KERNEL-VERSION   KUBELET-VERSION   CONTAINER-RUNTIME   HAS-UPDATES   HAS-DISRUPTIVE-UPDATES
master01  SUSE Linux Enterprise Server 15 SP1  4.12.14-197.10-default  v1.15.0        cri-o://1.15.0    <none>      <none>
worker10  SUSE Linux Enterprise Server 15 SP1  4.12.14-197.10-default  v1.15.0        cri-o://1.15.0    <none>      <none>
worker11  SUSE Linux Enterprise Server 15 SP1  4.12.14-197.10-default  v1.15.0        cri-o://1.15.0    <none>      <none>
tux@workstation:~/my-cluster>
```

---

## Summary:

In this exercise, you installed Skuba and the Kubernetes Client Utilities. You also tested that the SSH keys were set up properly and you could seamlessly connect to all of the nodes.

(End of Exercise)

## 2 Connecting to the Cluster

---

### Description:

In this section you will learn how to connect to the Cluster via both the command line and a GUI interface.

## 2- 1 Use the kubectl Command to Display Info About the Cluster

---

### Description:

In this exercise, you use the **kubectl** command to display info about the Kubernetes cluster.

---

### Task 1: Configure workstation for kubectl command

1. **On the Workstation, open a terminal**
2. Create official folder for kubernetes config filename  
**mkdir ~/.kube**
3. **On the Workstation, open a terminal**
4. Create config file

**cp ~/my-cluster/admin.conf ~/.kube/config**

### Task 2: Use the kubectl Command

1. From a terminal, enter the following command to view the URL of the Kubernetes master:

**kubectl cluster-info**

You should see the URL of the Kubernetes master displayed.

2. Enter the following command to see a more detailed output of the status of the cluster:

**kubectl cluster-info dump**

You should see a json dump of the status of cluster.

3. Enter the following command to display a list of Kubernetes nodes:

**kubectl get nodes**

You should see the list of nodes, their node names, status, age and version displayed.

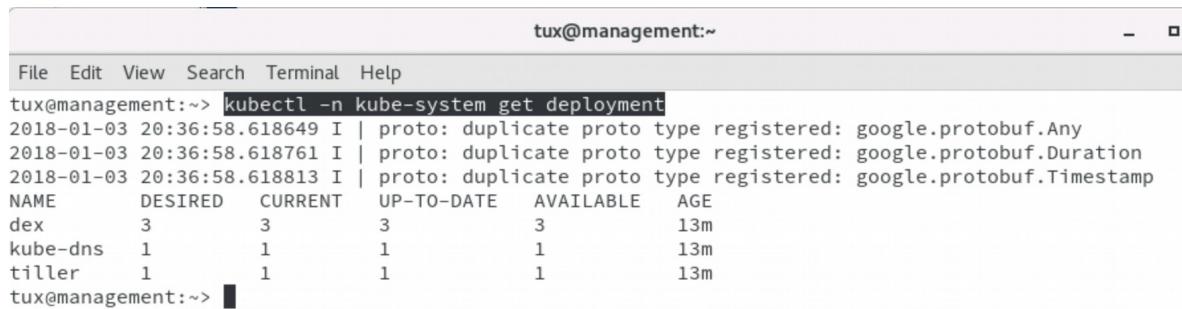
## SUSE U Advanced - SUSE CaaS Platform

```
tux@management:~/course_files/CAAS101/scripts> kubectl get nodes
2018-08-08 10:02:50.924418 I | proto: duplicate proto type registered: google.protobuf.Any
2018-08-08 10:02:50.924462 I | proto: duplicate proto type registered: google.protobuf.Duration
2018-08-08 10:02:50.924473 I | proto: duplicate proto type registered: google.protobuf.Timestamp
NAME      STATUS     AGE      VERSION
master01   Ready      1d      v1.9.8
master02   Ready      1d      v1.9.8
master03   Ready      1d      v1.9.8
worker10   Ready      1d      v1.9.8
worker11   Ready      1d      v1.9.8
worker12   Ready      1d      v1.9.8
tux@management:~/course_files/CAAS101/scripts> █
```

4. Enter the following command to display a list of Kubernetes nodes:

### **kubectl -n kube-system get deployments**

You should see the list of deployments, their nodes names, status, age and version displayed.



The screenshot shows a terminal window with the title "tux@management:~". The menu bar includes File, Edit, View, Search, Terminal, and Help. The command entered is "kubectl -n kube-system get deployment". The output shows deployment details for three components: dex, kube-dns, and tiller. The output is as follows:

```
tux@management:~> kubectl -n kube-system get deployment
2018-01-03 20:36:58.618649 I | proto: duplicate proto type registered: google.protobuf.Any
2018-01-03 20:36:58.618761 I | proto: duplicate proto type registered: google.protobuf.Duration
2018-01-03 20:36:58.618813 I | proto: duplicate proto type registered: google.protobuf.Timestamp
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
dex       3          3          3           3           13m
kube-dns  1          1          1           1           13m
tiller    1          1          1           1           13m
tux@management:~> █
```

5. Enter the following command to display a list of Kubernetes pods:

**kubectl -n kube-system get pods**

You should see the list of pods, their names, status, and age displayed.

```
tux@management:~$ kubectl -n kube-system get pods
2018-01-03 20:34:37.713518 I | proto: duplicate proto type registered: google.protobuf.Any
2018-01-03 20:34:37.713579 I | proto: duplicate proto type registered: google.protobuf.Duration
2018-01-03 20:34:37.713607 I | proto: duplicate proto type registered: google.protobuf.Timestamp
NAME                           READY   STATUS    RESTARTS   AGE
dex-3532614258-cssdz          1/1     Running   2          10m
dex-3532614258-nzjvq          1/1     Running   2          10m
dex-3532614258-xvpc9          1/1     Running   2          10m
haproxy-6c4cbef4e0545fc97a891c78b2f2b1c.infra.caasp.local 1/1     Running   0          9m
haproxy-b099b425f0194d18902aa68729c2a12e.infra.caasp.local 1/1     Running   0          10m
haproxy-d115a10918394ac88ab8c88bf7350556.infra.caasp.local 1/1     Running   0          10m
haproxy-dda6d06e26214c6d9c208ec2899748d2.infra.caasp.local 1/1     Running   0          10m
haproxy-e8185390bbc544efba4965e117a960eb.infra.caasp.local 1/1     Running   0          10m
haproxy-f7f84cefb02e4eadbcd03ce48642eef.infra.caasp.local 1/1     Running   0          9m
kube-dns-1144198277-3zf34      3/3     Running   0          11m
tiller-3881891813-4251j        1/1     Running   0          10m
tux@management:~$
```

---

**Summary:**

In this exercise, you downloaded the kubectl config file from the cluster and then ran some kubectl commands to view information about the cluster.

(End of Exercise)

## 2- 2 Deploy the Kubernetes Dashboard

---

### Description:

In this exercise, you deploy the Kubernetes Dashboard.

We will also apply an RBAC configuration that turns off token based authentication and stops the dashboard from timing out after 10 mins.

---

### Task 1: Deploy the Kubernetes Dashboard

1. On the **Workstation**, open a terminal
2. Enter the following command to deploy the Kubernetes Dashboard:

**kubectl apply -f ~/STW-CaaSPv4/manifests/dashboard**

You should see that a **secret**, **serviceaccount**, **role**, **rolebinding**, **deployment** and **service** were created:

3. Enter the following command to view the pods deployed in the **kube-system** namespace:

**kubectl -n kube-system get pods**

You should see the kubernetes-dashboard listed among the pods.

4. Enter the following command to view the deployments in the **kube-system** namespace:

**kubectl -n kube-system get deployments**

You should see the kubernetes-dashboard listed among the deployments.

5. Enter the following command to view the services in the **kube-system** namespace:

**kubectl -n kube-system get services**

You should see the kubernetes-dashboard listed among the services.

## Task 2: Access the Kubernetes Dashboard

1. Enter the following command to start a kubectl proxy in a screen session:

**screen kubectl proxy**

You should see the proxy session started. Notice that port 8001 on the local host is the port used.

2. Enter the following keystrokes to detach from the screen session:

**ctrl+a ctrl+d**

You should be detached from the screen session and back at a command prompt.

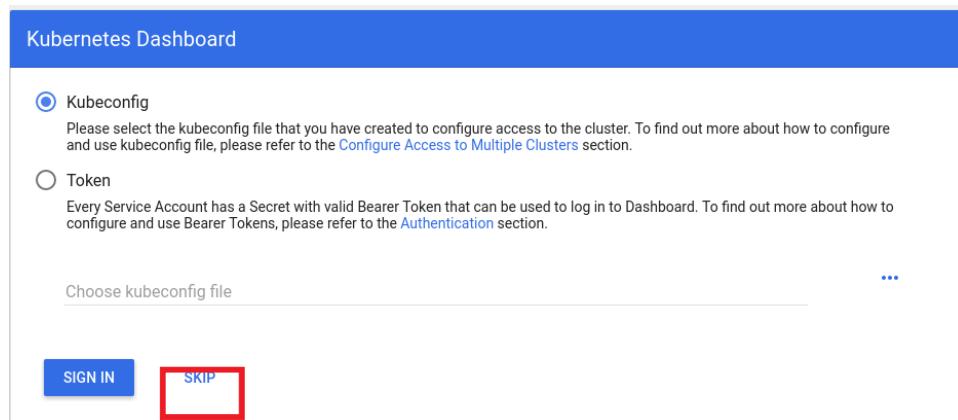
3. Open a web browser and point to:

**http://localhost:8001/api/v1/namespaces/kube-system/services/**

**https:kubernetes-dashboard:/proxy/#!/login**

You should see the Kubernetes Dashboard authentication method screen.

4. Click: **Skip** to login without a token



5. From the list on the left, in the **Namespace** section, from the drop-down list (which is probably displaying the **default** namespace), select **kube-system**  
You should see the kubernetes dashboard deployment in the **kube-system** namespace.
6. From that same drop-down menu, select **default**  
You should see that there is nothing deployed in the **default** namespace.

7.

---

**Summary:**

In this exercise, you deployed the Kubernetes Dashboard to the cluster. You then started a kubectl proxy in a screen session and accessed the Kubernetes Dashboard.

(End of Exercise)

## 2- 3 Install and Configure Helm

---

### Description:

In this exercise, you install Helm and config it to pull charts from SUSE.

---

### Task 1: Install Helm on Workstation

1. On the **Workstation**, open a terminal
2. Enter the following command to install skuba and the kubernetes client

**sudo zypper in helm**

It will tell you the amount of additional drive space needed and it will prompt you to press **y** to continue

### Task 2: Create Service account and Cluster Role for Helm

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to create a service account for Helm to use  
**kubectl create serviceaccount --namespace kube-system tiller**
3. Enter the following command (all one line) in the terminal to create a cluster role for Helm to use  
**kubectl create clusterrolebinding tiller --clusterrole=cluster-admin \ --serviceaccount=kube-system:tiller**

### Task 3: Initialize Helm

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to create a service account for Helm to use  
**helm init --tiller-image registry.suse.com/caasp/v4/helm-tiller:2.8.2 \ --service-account tiller**

3. Look to see if tiller is currently running (give it a minute if it is not in the **Ready** state)

### **kubectl get pods –all-namespaces**

```
tux@localhost:~> kubectl get pods --all-namespaces
NAMESPACE     NAME                               READY   STATUS    RESTARTS   AGE
kube-system   cilium-82kz9                         1/1    Running   0          47m
kube-system   cilium-bzsp7                         1/1    Running   0          47m
kube-system   cilium-operator-7d6dddbf5-qnbp5      1/1    Running   0          50m
kube-system   cilium-rwwjw                         1/1    Running   0          47m
kube-system   coredns-69c4947958-75vdp            1/1    Running   1          50m
kube-system   coredns-69c4947958-vmf6d            1/1    Running   3          50m
kube-system   etcd-master01                        1/1    Running   0          49m
kube-system   kube-apiserver-master01             1/1    Running   0          49m
kube-system   kube-controller-manager-master01      1/1    Running   0          49m
kube-system   kube-proxy-2nhh6                      1/1    Running   0          50m
kube-system   kube-proxy-ftbck                     1/1    Running   0          47m
kube-system   kube-proxy-pj698                      1/1    Running   0          48m
kube-system   kube-scheduler-master01              1/1    Running   0          49m
kube-system   kubernetes-dashboard-5f9c6b756f-hkwxh 1/1    Running   0          45m
kube-system   kured-4qzsf                          1/1    Running   0          45m
kube-system   kured-7tcfj                          1/1    Running   0          49m
kube-system   kured-9z4g7                          1/1    Running   0          47m
kube-system   oidc-dex-55fc689dc-mccw8            1/1    Running   1          50m
kube-system   oidc-gangway-7b7fbbdbdf-4w6z8        1/1    Running   0          50m
kube-system   tiller-deploy-7c666b7c99-7whpp       1/1    Running   0          42m
tux@localhost:~>
```

## Task 4: Add SUSE Charts to Helm

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to create a service account for Helm to use  
**helm repo add suse-charts <https://kubernetes-charts.suse.com>**
3. Look at available charts from SUSE

### **helm search suse**

```
tux@localhost:~> helm search suse
NAME          CHART VERSION APP VERSION DESCRIPTION
suse-charts/cf           2.17.1      1.4.1      A Helm chart for SUSE Cloud Foundry
suse-charts/cf-usb-sidecar-mysql 1.0.1       ...
suse-charts/cf-usb-sidecar-postgres 1.0.1       ...
suse-charts/console        2.4.0       2.4.0      A Helm chart for deploying Stratos UI Console
suse-charts/log-agent-rsyslog 1.0.1       8.39.0     Log Agent for forwarding logs of K8s control pl...
suse-charts/metrics         1.0.0       1.0.0      A Helm chart for Stratos Metrics
suse-charts/minibroker      0.2.0       ...
suse-charts/nginx-ingress   0.28.4      0.15.0     An nginx Ingress controller that uses ConfigMap...
suse-charts/uaa              2.17.1      1.4.1      A Helm chart for SUSE UAA
tux@localhost:~>
```

## Task 5: Install Centralize Logging

1. On the **Workstation**, open a terminal
2. Enter the following command in the terminal to create a service account for Helm to use

```
helm install suse-charts/log-agent-rsyslog --name 1.0.1 /  
--set server.host=rsyslog-server.default.svc.cluster.local /  
--set server.port=514
```

---

### Summary:

In this exercise, you installed Helm on the Management workstation. You then setup and configured it to pull charts directly from SUSE. You then installed the CaaS Platform's Centralized Logging Service.

(End of Exercise)

## 3 Deploying a Workload

---

### Description:

In this section you will deploy a simple workload in Kubernetes.

## 3- 1 Deploy a Simple Pod on Kubernetes

---

### Description:

In this exercise, you deploy the Nginx web server as a simple pod on the Kubernetes cluster.

---

### Task 1: View a Manifest for the Deployment

1. On the management workstation, in the text editor of your choice, open the file:

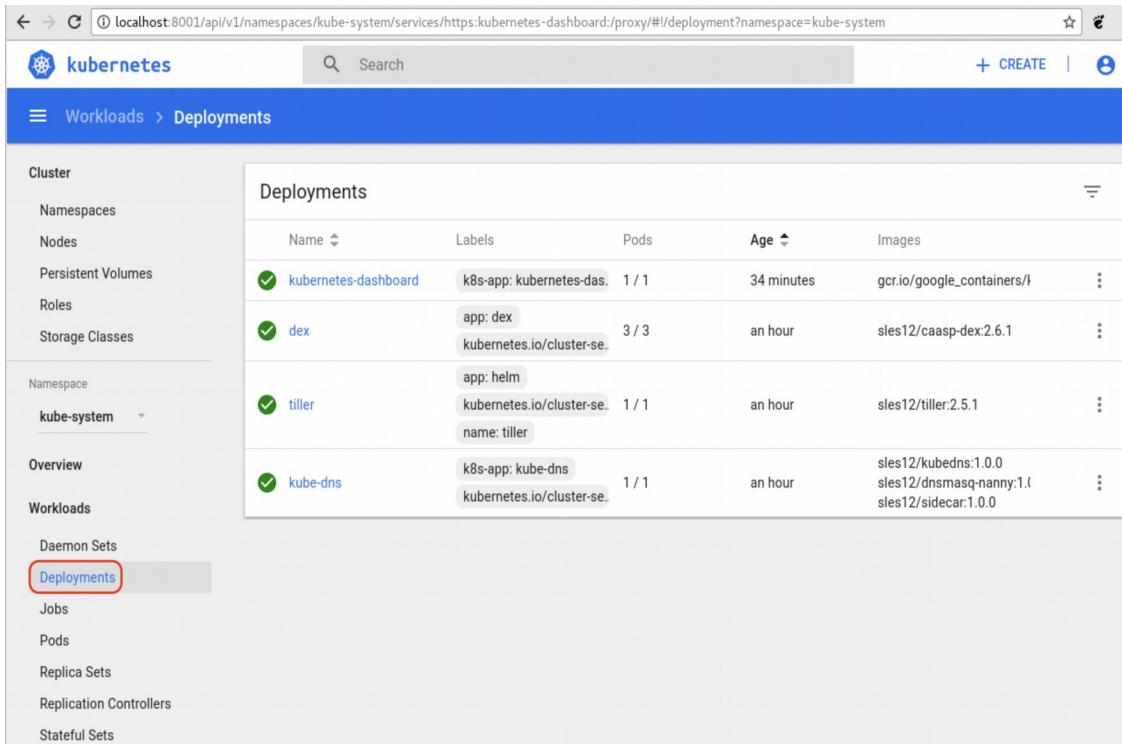
**~/STW-CaaSPv4/labs/nginx-deployment.yaml**

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 4
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: rmt.example.com:5000/sles12sp3_nginx
      ports:
        - containerPort: 80
```

## Task 2: Deploy a simple application (GUI option)

To deploy the application, go to the Kubernetes Dashboard

1. Click on **Deployments** on the left hand panel

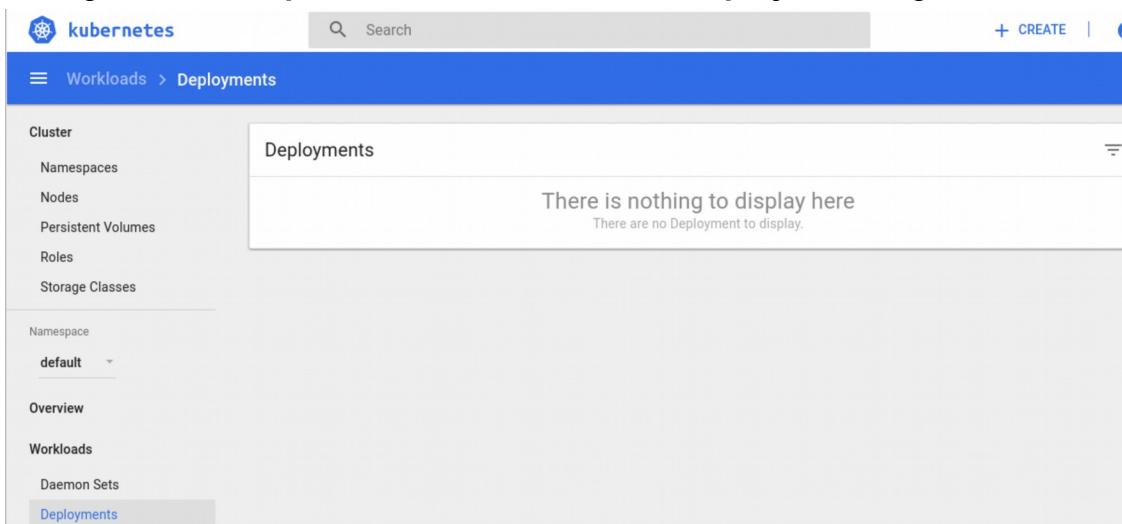


The screenshot shows the Kubernetes Dashboard interface. The left sidebar is titled 'Workloads' and includes options for Cluster, Namespaces, Nodes, Persistent Volumes, Roles, Storage Classes, Namespace (set to 'kube-system'), Overview, Workloads (Daemons Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets), and a 'Deployments' button which is highlighted with a red box. The main content area is titled 'Deployments' and lists four entries:

Name	Labels	Pods	Age	Images
kubernetes-dashboard	k8s-app: kubernetes-dash...	1 / 1	34 minutes	gcr.io/google_containers/...
dex	app: dex kubernetes.io/cluster-se...	3 / 3	an hour	sles12/caasp-dex:2.6.1
tiller	app: helm kubernetes.io/cluster-se...	1 / 1	an hour	sles12/tiller:2.5.1
kube-dns	k8s-app: kube-dns kubernetes.io/cluster-se...	1 / 1	an hour	sles12/kubedns:1.0.0 sles12/dnsmasq-nanny:1.0 sles12/sidecar:1.0.0

All of the currently deployed Deployments are in the kube-system namespace

2. Select **Namespace** on the left hand panel
3. Change the **Namespace to Default** and select **Deployments** again



The screenshot shows the Kubernetes Dashboard interface with the 'default' namespace selected in the 'Namespace' dropdown on the left sidebar. The 'Deployments' button is also highlighted with a red box. The main content area displays a message: 'There is nothing to display here' and 'There are no Deployment to display.'

note there are no current deployments in the **default** namespace

4. Click on **+ CREATE**
5. Select **Create from File** and the press ‘...’ to select file

CREATE FROM TEXT INPUT      **CREATE FROM FILE**      CREATE AN APP

---

Select YAML or JSON file specifying the resources to deploy to the currently selected namespace. [Learn more](#) 

Choose YAML or JSON file  

**UPLOAD**      [CANCEL](#)

6. Browse to **~/STW-CaaSPv4/labs/nginx-deployment.yaml**
7. Press Upload
8. Select **Deployments** view and we should now see the nginx-deployment

 **kubernetes**       Search      + CREATE | 

☰ Workloads > Deployments

Name	Labels	Pods	Age	Images
nginx-deployment	app: nginx	1 / 1	32 seconds	nginx:1.7.9

Cluster  
Namespaces  
Nodes  
Persistent Volumes  
Roles  
Storage Classes

Namespace  
default

Overview  
Workloads  
Daemon Sets  
**Deployments**

## 9. Explore the Kubernetes Dashboard

Make sure you look at the **Overview, Workloads, and Pods** views

The screenshot shows the Kubernetes Dashboard interface. On the left, there's a sidebar with navigation links: Cluster, Namespaces, Nodes, Persistent Volumes, Roles, Storage Classes, Namespace (default), Overview, Workloads (which is selected and highlighted in blue), Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, and Stateful Sets. The main content area has three tabs: Deployments, Pods, and Replica Sets. The Deployments tab shows one deployment named "nginx-deployment" with 4 pods, all running. The Pods tab shows four pods corresponding to the deployment, each running on a different worker node. The Replica Sets tab shows one replica set with 4 pods, all running.

## Task 3: Deploy a simple application (Command Line option)

10. To deploy the pod, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-deployment.yaml**

```
tux@management:~> kubectl apply -f ~/course_files/CAAS101/manifests/labs/nginx-deployment.yaml
2018-03-13 09:27:42.158048 I | proto: duplicate proto type registered: google.protobuf.Any
2018-03-13 09:27:42.158121 I | proto: duplicate proto type registered: google.protobuf.Duration
2018-03-13 09:27:42.158136 I | proto: duplicate proto type registered: google.protobuf.Timestamp
deployment "nginx-deployment" created
tux@management:~>
```

You should see the deployment “nginx-deployment” was created.

11. Enter the following command to view the deployments:

**kubectl get deployments**

You should see that a single instance of the `nginx-deployment` pod is running.

12. Enter the following command to view the deployed pods:

**kubectl get pods**

You should see a single instance of the `nginx-deployment` pod running.

---

**Summary:**

In this exercise, you launched a single instance of the Nginx web server as a pod in a deployment on the cluster.

(End of Exercise)

## 3- 2 Delete a Deployment on Kubernetes

### Description:

In this exercise, you delete the Nginx web server deployment that was previously deployed on the Kubernetes cluster.

### Task 1: Delete the Deployment (GUI Option)

1. Select **nginx-deployment** under Deployments
2. Select **Delete**

The screenshot shows the Kubernetes UI for managing workloads. In the top navigation bar, 'Deployments' is selected under 'Workloads'. The main panel displays the 'nginx-deployment' details, including its name, namespace, labels, selector, strategy, and current status (10 updated, 10 total, 10 available, 0 unavailable). Below this, a 'New Replica Set' table lists one pod named 'nginx-deployment-431080' with 10/10 replicas, created 39 minutes ago, and using the 'nginx:1.7.9' image. The 'DELETE' button in the top right corner of the main panel is highlighted with a red box.

### Task 2: Delete the Deployment (Command Line option)

3. To view the deployments, on the **Workstation**, enter the following command:

**kubectl get deployments**

You should see that one instance of the **nginx-deployment** is running.

4. Enter the following command to delete the deployment:

**kubectl delete deployment nginx-deployment**

You should see the **nginx-deployment** was deleted.

5. View the deployments again:

**kubectl get deployments**

You should see that the `nginx-deployment` is no longer running.

6. View the pods:

**kubectl get pods**

You should see that all of the pods that were part of the `nginx-deployment` are no longer running.

---

**Summary:**

In this exercise, you deleted the Nginx web server deployment that was previously deployed on the cluster.

(End of Exercise)

## 3- 3 Scale Out a Deployment

---

### Description:

In this exercise, you scale out a running Deployment.

---

### Task 1: View the Manifest for the Deployment

1. In the text editor of your choice, view the  
**~/STW-CaaSPv4/labs/nginx-deployment.yaml** file

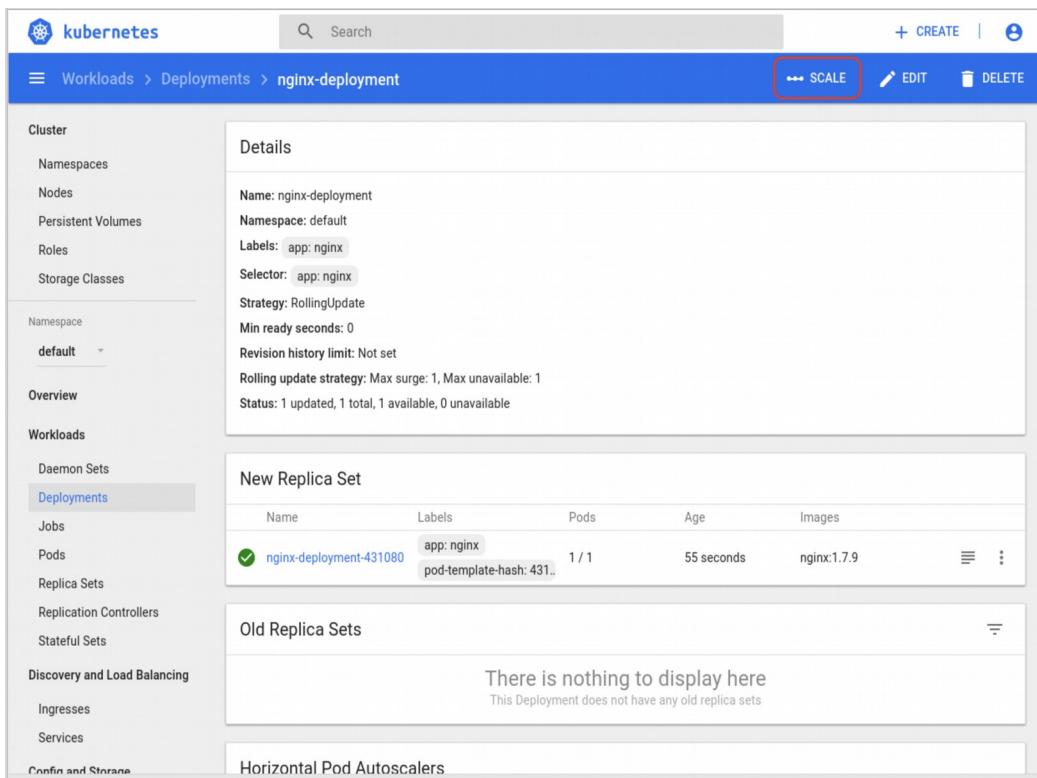
```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 4
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: rmt.example.com:5000/sles12sp3_nginx
          ports:
            - containerPort: 80
```

### Task 2: Scale the Deployment (GUI option)

1. From **Kubernetes Dashboard** select **Create** and open  
**~/STW-CaaSPv4/labs/nginx-scale.yaml**
2. Press the **Deploy** button
3. Select the **nginx-deployment** under **Deployments**

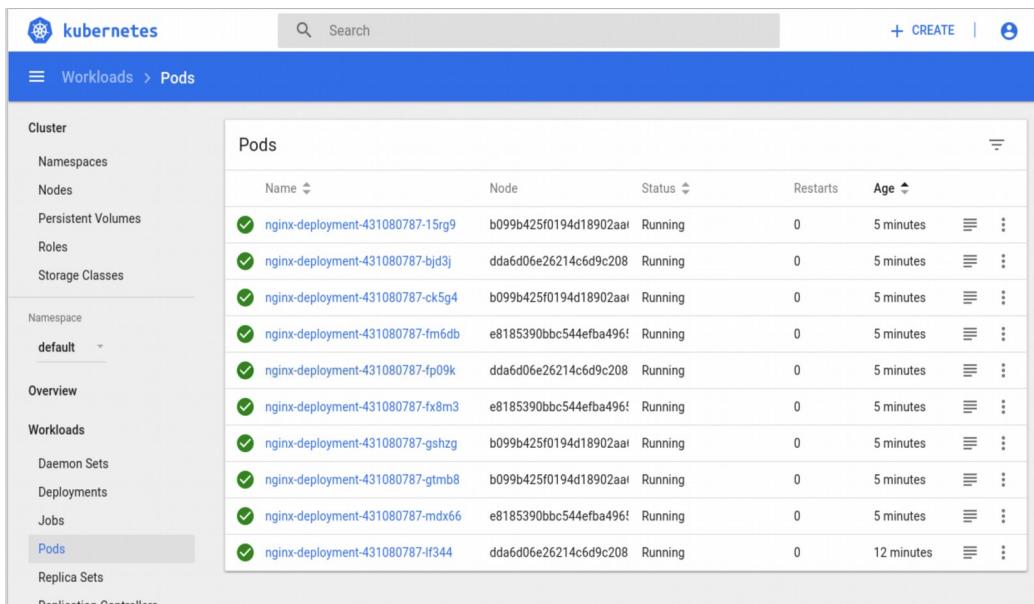
## SUSE U Advanced - SUSE CaaS Platform

4. click the **Scale** Button and change the number of **Desired pods** to 10



The screenshot shows the Kubernetes UI for managing a Deployment named 'nginx-deployment'. The 'Deployment' tab is selected in the sidebar. The 'SCALE' button is highlighted with a red box. The 'Details' section shows the deployment's configuration: Name: nginx-deployment, Namespace: default, Labels: app: nginx, Selector: app: nginx, Strategy: RollingUpdate, Min ready seconds: 0, Revision history limit: Not set, Rolling update strategy: Max surge: 1, Max unavailable: 1, Status: 1 updated, 1 total, 1 available, 0 unavailable. Below this, the 'New Replica Set' table lists one replica set ('nginx-deployment-431080') with 1 pod. The 'Old Replica Sets' section is empty. At the bottom, there is a 'Horizontal Pod Autoscalers' section.

5. View the deployed pods by selecting **Pods** in the left hand panel



The screenshot shows the Kubernetes UI for viewing pods. The 'Pods' tab is selected in the sidebar. The main table displays multiple pods from the 'nginx-deployment' deployment, all in a 'Running' state. The columns include Name, Node, Status, Restarts, and Age. The pods listed are: nginx-deployment-431080787-15rg9, nginx-deployment-431080787-bjd3j, nginx-deployment-431080787-ck5g4, nginx-deployment-431080787-fm6db, nginx-deployment-431080787-fp09k, nginx-deployment-431080787-fx8m3, nginx-deployment-431080787-gshzg, nginx-deployment-431080787-gtmb8, nginx-deployment-431080787-mdx66, and nginx-deployment-431080787-if344.

6. Scale the Pods back down to 4

## Task 3: Scale the Deployment (Command Line Option)

1. In the text editor of your choice, open the **~/STW-CaaSPv4/labs/nginx-scale.yaml** file

Note the the only difference is the number of replicas

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 2
  revisionHistoryLimit: 5
  minReadySeconds: 20
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 25%
      maxSurge: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: rmt.example.com:5000/sles12sp3_nginx
          ports:
            - containerPort: 80
```

2. To scale the deployment, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-scale.yaml**

3. Enter the following command to view the deployments:

**kubectl get deployments**

You should see that one instance of the **nginx-deployment** Deployment is running.

4. Enter the following command to view the deployed pods:

**kubectl get pods**

You should see 10 instances of the **nginx-deployment** pod running.

---

### **Summary:**

In this exercise, you created a new manifest for an existing deployment that specified a smaller number of replicas. You then applied the updated manifest to scale in the deployment. Finally you applied the original manifest to scale the deployment back out.

(End of Exercise)

## **4 Working With Kubernetes**

---

### **Description:**

In this section you are introduced to how to work with Kubernetes. You are first introduced to Kubernetes configuration and management utilities and manifests. You then deploy and manage pods on a Kubernetes cluster.

### **4- 1 Update a Deployment**

---

#### **Description:**

In this exercise, you update a running pod.

---

#### **Task 1: Create a New Manifest for the Deployment**

1. In the text editor of your choice, open the **~/STW-CaaSPv4/labs/nginx-update.yaml** file
2. Notice the section in **red**

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 4
  template:
    metadata:
```

**labels:**

**app: nginx**

**spec:**

**containers:**

- **name: nginx**

**image: rmt.example.com:5000/nginx:1.9.0**

**ports:**

- **containerPort: 80**

## Task 2: Update the Deployment

1. If you don't already have **nginx** deployed Deploy either through **KubeDashboard** using the file **~/STW-CaaSPv4/labs/nginx-pre-update.yaml** or the following command line:

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-pre-update.yaml**

2. To display information on the current nginx deployment enter the following command:

**kubectl describe deployment -l app=nginx**

You should see the description of the nginx deployment displayed.

Notice the image version is: nginx:1.7.9

3. Open another terminal and enter the following command to watch the running pods:

**watch kubectl get pods**

You should see a list of the running pods displayed with the list updating every 2 seconds.

4. To update the deployment, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-update.yaml**

You should see the deployment "nginx-deployment" was configured.

5. Enter the following command to view the deployments:

**kubectl get deployments**

You should see that 4 instances of the **nginx-deployment** pod are DESIRED and, depending on when you ran the command, the values in the CURRENT, UP-TO-DATE and AVAILABLE columns may be more, fewer or the same number as the update happens .

6. In the terminal where you are watching the pods enter **ctrl+c** to stop the watch command
7. Enter the following command to display information about the running deployment of nginx:

**kubectl describe pods -l app=nginx**

Notice the image version is now: nginx:1.9.0

---

**Summary:**

In this exercise, you created a new manifest to update the running nginx deployment. You then updated the deployment and verified that it was updated.

(End of Exercise)

## 4- 2 Update a Deployment Via Rolling Updates

---

### Description:

In this exercise, you update a running pod using rolling updates.

---

### Task 1: Create a New Manifest for the Deployment

1. On the management workstation, In the text editor of your choice, open the **~/STW-CaaSPv4/labs/nginx-rolling\_update.yaml** file
2. Notice the changes are in **red**

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 4
  revisionHistoryLimit: 5
  minReadySeconds: 20
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 25%
      maxSurge: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: rmt.example.com:5000/nginx:1.12.0
          ports:
            - containerPort: 80
```

3. Save the file and close the text editor

## Task 2: Update the Deployment

1. To display information on the current nginx deployment enter the following command:

**kubectl describe deployment -l app=nginx**

You should see the description of the nginx deployment displayed.

Notice the image version is: `nginx:1.9.0`

2. Open another terminal and enter the following command to watch the running pods:

**watch kubectl get pods**

You should see a list of the running pods displayed with the list updating every 2 seconds.

3. To update the deployment, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-rolling\_update.yaml**

You should see the deployment “nginx-deployment” was configured.

In the terminal where you are watching the pods you should see the number of running pods drop to 3 (because of maxUnavailable: 25%) and 3 new pods start deploying. Shortly after the 3 new pods are running you should see the number of running pods scale up to 6 (because of maxSurge: 2) and then back to 4.

4. Enter the following command to view the deployments:

**kubectl get deployments**

You should see that 4 instances of the `nginx-deployment` pod are running.

5. In the terminal where you are watching the pods enter **ctrl+c** to stop the watch command
6. Enter the following command to display information about the running deployment of nginx:

**kubectl describe deployment -l app=nginx**

Notice the image version is now: `nginx:1.12.0`

---

### Summary:

In this exercise, you created a new manifest to update the running nginx deployment using the rolling update type. You then updated the deployment and verified that it was updated.

(End of Exercise)

---

## 4- 3 Expose a Service Running in a Pod

---

### Description:

In this exercise, you expose a service running in a pod.

---

### Task 1: Create a Manifest for the Service

1. On the Workstation, in the text editor of your choice, open the file:

**~/STW-CaaSPv4/labs/nginx-service.yaml**

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 30000
  selector:
    app: nginx
```

### Task 2: Define the Service

1. To define the service in the cluster, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-service.yaml**

You should see the service “nginx-service” was created.

2. Enter the following command to view the services:

**kubectl get services**

You should see that the **nginx-service** service is defined. Notice the 80:30000 under ports showing external port 30000 will be redirected into internal port 80.

### Task 3: Access the Exposed Service

1. To access the exposed service, open a web browser and point to:

**http://worker.example.com:30000**

You should see the Welcome to nginx web page.



**Note:**

If desired, you can change the URL to point to a specific worker node (i.e. **worker10.example.com**) and see that the web page is still accessible. This demonstrates that the kube-proxy is working on all of the cluster nodes.

---

**Summary:**

In this exercise, you defined a service in the cluster exposing port 80 that allowed access to the nginx pod running on the cluster. You then accessed the nginx pod in a web browser.

(End of Exercise)

---

## 4- 4 Setup Readiness and Liveness Probes

---

**Description:**

In this exercise, you learn how to setup and configure Readiness and Liveness Probes

---

### Task 1: View manifest for the Deployment

1. In the text editor of your choice, open the **~/STW-CaaSPv4/labs/nginx-deployment-health.yaml** file

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 4
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: rmt.example.com:5000/sles12sp3_nginx
        ports:
        - containerPort: 80
      readinessProbe:
        httpGet:
          path: /
          port: 80
        initialDelaySeconds: 5
        timeoutSeconds: 1
        periodSeconds: 15
      livenessProbe:
        httpGet:
          path: /
          port: 80
        initialDelaySeconds: 15
        timeoutSeconds: 1
        periodSeconds: 15
```

2. Notice the ReadinessProbe and LivenessProbe settings

## Task 2: Deploy Manifest

1. Deploy the manifest

**kubectl apply -f ~/STW-CaaSPv4/labs/nginx-deployment-health.yaml**

## Task 3: View the Replica sets

2. On the lefthand side of the Kubernetes Dashboard, under Workload, select Replica Sets

Name	Node	Status	Restarts	Age	CPU (cores)	Memory (bytes)
nginx-deploymer	worker12	Running	1	an hour	0	9.359 Mi
nginx-deploymer	worker11	Running	0	an hour	0	9.504 Mi
nginx-deploymer	worker10	Running	1	an hour	0	9.348 Mi
nginx-deploymer	worker10	Running	0	an hour	0	9.355 Mi

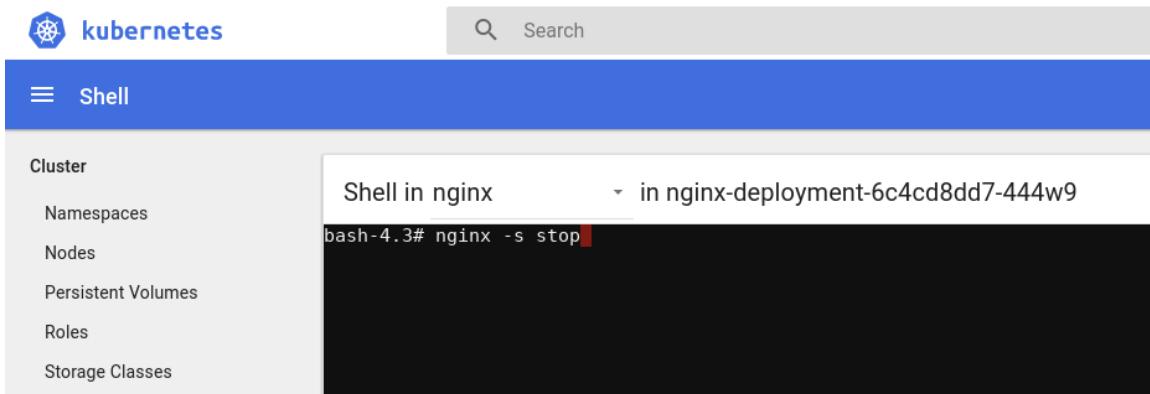
3. Notice the number of Restarts on the container

## Task 4: Kill an nginx server

1. Select one of the pods labeled nginx-deployment-#### by clicking on it
2. Click the EXEC button to be taken to a shell prompt inside the running pod

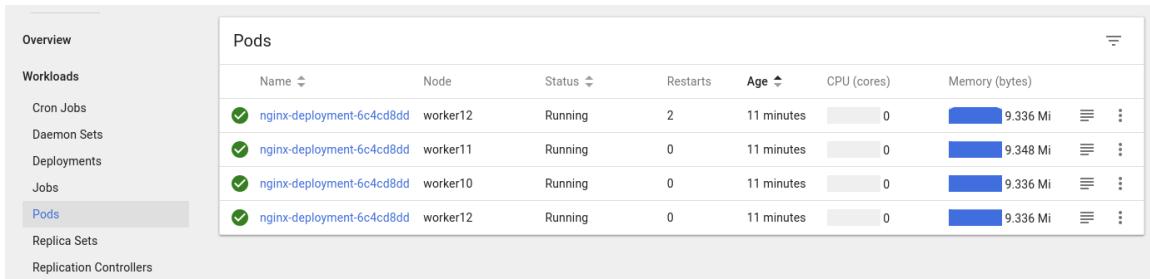
Cluster	CPU usage	Memory usage ⓘ
Namespaces		

- From the Shell enter the following command to kill the nginx process  
**nginx -s stop**



The screenshot shows the Kubernetes UI with a "Shell" tab selected. On the left, there's a sidebar with "Cluster" navigation options: Namespaces, Nodes, Persistent Volumes, Roles, and Storage Classes. The main area is titled "Shell in nginx" and "in nginx-deployment-6c4cd8dd7-444w9". It displays a terminal session with the command "bash-4.3# nginx -s stop" entered.

- Click on Pods to view the current pods running



The screenshot shows the Kubernetes UI with the "Pods" tab selected under the "Workloads" section. The sidebar has "Overview" and "Replica Sets" sections. The main table lists four pods from the "nginx-deployment" deployment, each running on a different worker node (worker12, worker11, worker10, worker12) and having 0 restarts.

Name	Node	Status	Restarts	Age	CPU (cores)	Memory (bytes)
nginx-deployment-6c4cd8dd	worker12	Running	2	11 minutes	0	9.336 Mi
nginx-deployment-6c4cd8dd	worker11	Running	0	11 minutes	0	9.348 Mi
nginx-deployment-6c4cd8dd	worker10	Running	0	11 minutes	0	9.336 Mi
nginx-deployment-6c4cd8dd	worker12	Running	0	11 minutes	0	9.336 Mi

- Notice the number of restarts has increased. This is because Kubernetes could no longer reach the server on port 80 so it restarted that Nginx Container inside the Pod

---

### Summary:

In this exercise, you created a new manifest to update the running nginx deployment. You then updated the deployment and verified that it was updated.

(End of Exercise)

## 4- 5 Define Limits for Containers and Pods in Kubernetes

---

### Description:

In this exercise, you define limits for containers and pods in the Kubernetes cluster.

---

## Task 1: Create a New Namespace in the Cluster

1. On the Workstation, at the command line, enter the following command to create a new namespace in the Kubernetes cluster:

**kubectl create namespace limit-example**

You should see that a new namespace was created.

2. Enter the following command to display the namespaces:

**kubectl get namespaces**

You should see the new namespace listed.

## Task 2: Create a Manifest for the Limits

1. In the text editor open the file:

**~/STW-CaaSPv4/labs/limits.yaml**

```
apiVersion: v1
kind: LimitRange
metadata:
  name: mylimits
spec:
  limits:
  - max:
      cpu: "2"
      memory: 1Gi
    min:
      cpu: 200m
      memory: 6Mi
    type: Pod
  - default:
      cpu: 300m
      memory: 200Mi
    defaultRequest:
      cpu: 200m
      memory: 100Mi
    max:
      cpu: "2"
      memory: 1Gi
    min:
      cpu: 100m
      memory: 3Mi
    type: Container
```

## Task 3: Apply the Limits to the Namespace

1. To deploy the pod, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/limits.yaml --namespace=limit-example**

You should see the limitrange “mylimits” was created.

2. Enter the following command to display the limitranges:

**kubectl describe limitranges --namespace=limit-example**

You should see the **mylimits** limitrange listed.

## SUSE U Advanced - SUSE CaaS Platform

Name:	mylimits	Namespace:	limit-example	Type	Resource	Min	Max	Default Request	Default Limit	Max Limit/Request Ratio
Pod	cpu			Pod	cpu	200m	2	-	-	-
Pod	memory			Pod	memory	6Mi	1Gi	-	-	-
Container	cpu			Container	cpu	100m	2	200m	300m	-
Container	memory			Container	memory	3Mi	1Gi	100Mi	200Mi	-

---

### **Summary:**

In this exercise, you created a new namespace in the Kubernetes cluster. You then defined limits for containers and pods and applied them to the new namespace.

(End of Exercise)

---

## **4- 6 Introduction to Helm**

---

### **Description:**

In this exercise, you are introduced to the helm utility.

---

### **Task 1: Use Some Basic Helm Commands**

1. In a terminal, enter the following command to source the Helm auto-completion functions into your shell environment:

**source <(helm completion bash)**

Your shell environment should now have the helm auto-completion functions available.

2. Enter the following command to display a list of available helm charts:

**helm search**

You should see a list of available helm charts listed.

3. Try running the command again but rather than typing in the entire command, only type:

**helm sea[Tab]**

(Where [Tab] is the tab key)

Notice that the “**search**” option is auto-completed.

4. Enter the following command to search for a specific helm chart:

**helm search dokewiki**

You should see the dokewiki chart listed.

5. Enter the following command to display the list of currently configured repos:

### **helm repo list**

You should see a repo for <https://kubernetes-charts.storage.googleapis.com>, [kubernetes-charts.suse.com](https://kubernetes-charts.suse.com) and an addition to the local repo.

---

#### **Summary:**

In this exercise, you initialized helm. You then ran some basic helm commands such as configuring auto-completion, searching for helm charts and displaying repositories.

(End of Exercise)

## **4- 7 Deploy an Application with Helm**

---

#### **Description:**

In this exercise, you deploy an application from a helm chart using the helm command.

---

### **Task 1: Create Helm Chart Config File**

1. On the Workstation, in a terminal, enter the following command to search for a dokuwiki helm chart:

#### **helm search dokuwiki**

You should see a helm chart named stable/dokuwiki listed with its available chart version.

2. Enter the following command to view the default configuration for the dokuwiki chart:

#### **helm inspect stable/dokuwiki | less**

You should see the configuration displayed in the less pager. Page through the configuration to see what variables are being set.

3. In the text editor of your choice, create/open the **~/STW-CaaSPv4/labs/dokuwiki-config.yaml** file
4. Enter the following in the file:

**dokuwikiPassword: password123**  
**serviceType: NodePort**  
**persistence:**  
**enabled: false**

5. Save the file and close the text editor

## Task 2: Deploy the Helm Chart

1. In a terminal, enter the following command to view the current helm releases:

### **helm list**

You should not see the **dokuwiki** application listed.

2. Enter the following command to deploy the chart:

### **helm install -f ~/STW-CaaSPv4/labs/dokuwiki-config.yaml --name mywiki stable/dokuwiki**

You should see the chart was deployed.

In the output of the chart deployment, in the **==> v1/Service** section, notice the IP and port(s) the application is listening on. The NodePort(s) that the application is listening on are the number after the colon (**:**) in the **PORT(S)** column.

Example: **80:32313/TCP,443:31034/TCP**

In this example the NodePorts are **32313** for http and **31034** for https.

Record the http NodePort here:

**DOKUWIKI\_PORT=**\_\_\_\_\_

3. Enter the following command to display the current helm releases:

### **helm list**

You should now see the **dokuwiki** chart listed with a name of **mywiki**. Notice the **REVISION** column shows the number **1** as this is the initial deployment of this release.

4. Enter the following command to view the release history for the application:

### **helm history mywiki**

You should see that only a single release of **mywiki** exists.

5. Enter the following command to view the status of the **mywiki** release:

### **helm status mywiki**

You should see output similar to what was displayed when the chart was first deployed.

6. You can also enter the following **kubectl** commands:

**kubectl get deployments**

**kubectl get pods**

**kubectl get services**

Notice that you see that same info about the deployed deployments/pods/services as if you were to have deployed them from manifests.

## Task 3: Access the Application Deployed by the Helm Chart

1. Open a web browser and point to:

## **[http://worker.example.com:DOKUWIKI\\_PORT](http://worker.example.com:DOKUWIKI_PORT)**

You should see the **My Wiki** page displayed.

2. On the top right of the page click: **Login**
3. Enter the following credentials:

**Username:** user

**Password:** password123

You should be logged in.

After logging in you probably see a number of warnings of available hotfixes and/or updates.

## **Task 4: Update the Application Release**

1. On the **Workstation**, in a terminal, make a copy of the chart config file you created at the beginning of this exercise:

**cp ~/dokewiki-config.yaml ~/dokewiki-config-update.yaml**

2. In the text editor of your choice, open the **~/dokewiki-config-update.yaml** file
3. Edit the file to match the following (changes are in **red**):

**images: bitnami/dokewiki:latest**

**dokewikiPassword: password123**

**serviceType: NodePort**

**persistence:**

**enabled: false**

4. Save the file and close the text editor

5. Enter the following command to upgrade the **mywiki** release:

**helm upgrade -f ~/dokewiki-config-update.yaml mywiki stable/dokewiki**

You should see the chart was successfully deployed.

6. Enter the following command to view the current helm releases:

**helm list**

You should see the **dokewiki** chart named **mywiki** listed. Notice the **REVISION** column now shows the number **2** as the release has been updated once.

7. Enter the following command to display the release history for the **mywiki** release:

**helm history mywiki**

Notice that there are 2 revisions. Also notice the the first revision's **STATUS** is **SUPERSEDED** and the second revision's **STATUS** is **DEPLOYED**.

8. In the web browser, refresh the web page (or log back in if you have logged out)

You should no longer see the warning messages about available hotfixes and/or

updates.

## Task 5: Delete a Deployed Helm Chart Release

1. In a terminal, enter the following command to delete the **mywiki** release:

**helm delete mywiki**

You should see that the release was deleted.

2. Enter the following command to list the current helm releases:

**helm list**

You should no longer see the **dokewiki** chart named **mywiki** displayed.

3. Enter the following command to display the status of the **mywiki** release:

**helm status mywiki**

Notice that the **STATUS** is **DELETED** but also that it remembered that it had been deployed as it has a date listed in **LAST DEPLOYED**.

---

### Summary:

In this exercise, you first deployed a helm chart. You then accessed the application that was deployed. Next you updated the release, verifying it was updated. Finally you deleted the release.

(End of Exercise)

---

## 5 Storage

---

### Description:

How to handle persistent storage in pods.

## 5- 1 Configure NFS Persistent Storage

---

### Description:

In this exercise, you configure a persistent volume on an NFS server. You then create a pod that updates a file on the persistent volume and a pod that exports the file via http

---

### Task 1: Create the Persistent Volume on the NFS Server

1. On the RMT server, as the **root** user with the password of **Linux** enter the following commands in a terminal session to create the persistent volume and **index.html** file:

```
mkdir -p /export/vol-01
touch /export/vol-01/index.html
```

The directory and file should now exist.

### Task 2: Create the Manifest for the Persistent Volume

2. On the CAAS101-management VM, in the text editor of your choice, open the file:

**~/STW-CaaSPv4/labs/pv/nfs-pv-vol-01.yaml**

Enter the following in the file:

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs-vol-01
spec:
  capacity:
    storage: 1Mi
  accessModes:
    - ReadWriteMany
  nfs:
    # FIXME: use the right IP
    server: 192.168.110.2
    path: "/export/vol-01"
```

### Task 3: Create the Manifest for the Persistent Volume Claim

1. In the text editor of your choice, open the file:

### ~/STW-CaaSPv4/labs/pv/nfs-pvc-vol-01.yaml

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: nfs-vol-01
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 1Mi
```

### Task 4: Create the Manifest for the Busybox Instance

1. In the text editor of your choice, open the file:

### ~/STW-CaaSPv4/labs/pv/nfs-busybox-deployment.yaml

```
# This mounts the nfs volume claim into /mnt and continuously
# overwrites /mnt/index.html with the time and hostname of the pod.

apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nfs-busybox
spec:
  replicas: 1
  template:
    metadata:
      labels:
        name: nfs-busybox
    spec:
      containers:
        - image: smt.example.com:5000/busybox
          command:
            - sh
            - -c
            - 'while true; do date > /mnt/index.html; hostname >> /mnt/index.html; sleep $((RANDOM % 5 + 5)); done'
          imagePullPolicy: IfNotPresent
        name: busybox
      volumeMounts:
        # name must match the volume name below
        - name: nfs-vol-01
          mountPath: "/mnt"
    volumes:
      - name: nfs-vol-01
        persistentVolumeClaim:
          claimName: nfs-vol-01
```

### Task 5: Create the Manifests for the Web Frontend

1. In the text editor of your choice, create/open the file:

### ~/STW-CaaSPv4/labs/pv/nfs-web-deployment.yaml

```
# This pod mounts the nfs volume claim into /usr/share/nginx/html and
# serves a simple web page.

apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nfs-web
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: web-frontend
    spec:
      containers:
        - name: web
          image: rmt.example.com:5000/nginx:1.9.0
          ports:
            - name: web
              containerPort: 80
      volumeMounts:
        - name: nfs-vol-01
          mountPath: "/usr/share/nginx/html"
  volumes:
    - name: nfs-vol-01
      persistentVolumeClaim:
        claimName: nfs-vol-01
```

2. Next, open the file:

**~/STW-CaaSPv4/labs/pv/nfs-web-service.yaml**

```
apiVersion: v1
kind: Service
metadata:
  name: nfs-web
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 30080
  selector:
    app: web-frontend
```

## Task 6: Deploy the Objects

1. To deploy the volumes/pods/service, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/pv**

You should see the following were created (not necessarily in this order):

deployment “nfs-busybox”  
persistentvolume “nfs-vol-01”  
persistenvolumeclaim “nfs-vol-01”  
deployment “nfs-web”  
deployment “nfs-web”

2. Enter the following command to view the deployments:

### **kubectl get deployments**

You should see the **nfs-busybox** and **nfs-web** deployments listed.

3. Enter the following command to view the pods:

### **watch kubectl get pods**

You should see the pods for the **nfs-busybox** and **nfs-web** deployments listed.

4. Enter the following command to view the persistent volumes:

### **kubectl get pv**

You should see the persistent volume **nfs-vol-01** listed.

5. Enter the following command to view the persistent volumes:

### **kubectl get pvc**

You should see the persistent volume claim **nfs-vol-01** listed.

## Task 7: Test the Persistent Data

1. On the **SMT server**, enter the following command to view the content of the **index.html** file on the NFS server:

### **cat /export/vol-01/index.html**

You should see a line with the date and time and a line with the pod name of the **nfs-busybox** pod. If you rerun the command you should see that the time stamp is updating.

2. On the **management workstation**, open a web browser and point to:

**http://worker10.example.com:30080**

You should see the content of the **index.html** file. When you refresh the page you should see the time stamp updating. (Also notice that if you change the URL to the different worker nodes you will see the same thing.)

## Task 8: Remove the Objects from the Cluster

1. In the first terminal, enter the following commands to delete all of the objects:

```
kubectl delete service nfs-web
kubectl delete deployment nfs-web
kubectl delete deployment nfs-busybox
kubectl delete pv nfs-vol-01
kubectl delete pvc nfs-vol-01
```

You should see that the objects were deleted.

---

### Summary:

In this exercise, you created manifests for a persistent volume, a persistent volume claim, a pod to that attached to the volume and writes data to an index.html file in the volume, and a web server that attaches to the volume and displays the index.html file. You then verified that the index.html file was being updated by looking at the file both on the NFS volume and the web server.

(End of Exercise)

## 5- 2 Configure Persistent Storage with a NFS StorageClass

---

### Description:

In this exercise, you define a StorageClass that will automatically create volumes on an NFS server. You then create a volume claim that will cause the volume to be created on the NFS server and a pod that will write a file in the volume.

---

### Task 1: Create the Directory for the Persistent Volumes on the NFS Server

1. On the **SMT server**, as the **root** user with a password of **Linux**, enter the following commands to create the directory that will contain the volumes:

```
mkdir -p /export/volumes
chmod 777 /export/volumes
```

The directory and file should now exist.

### Task 2: Define a Service Account, Cluster Role and Cluster Role Binding

1. On the **CAAS101-management VM**, in the text editor of your choice, open the file:  
**~/STW-CaaSPv4/labs/nfs-client-storageclass/auth/serviceaccount.yaml**

```
apiVersion: v1
kind: Service
metadata:
  name: nfs-web
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 30080
  selector:
    app: web-frontend
```

2. In the text editor of your choice, open the file:
3. **~/STW-CaaSPv4/labs/nfs-client-storageclass/auth/clusterrole.yaml**

```
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: nfs-client-provisioner-runner
rules:
  - apiGroups: [""]
    resources: ["persistentvolumes"]
    verbs: ["get", "list", "watch", "create", "delete"]
  - apiGroups: [""]
    resources: ["persistentvolumeclaims"]
    verbs: ["get", "list", "watch", "update"]
  - apiGroups: ["storage.k8s.io"]
    resources: ["storageclasses"]
    verbs: ["get", "list", "watch"]
  - apiGroups: [""]
    resources: ["events"]
    verbs: ["list", "watch", "create", "update", "patch"]
```

4. In the text editor of your choice, open the file:

**~/STW-CaaSPv4/labs/nfs-client-storageclass/auth/  
clusterrolebinding.yaml**

```
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: run-nfs-client-provisioner
subjects:
  - kind: ServiceAccount
    name: nfs-client-provisioner
    namespace: default
roleRef:
  kind: ClusterRole
  name: nfs-client-provisioner-runner
  apiGroup: rbac.authorization.k8s.io
```

5. Enter the following command to open the service, clusterrole and clusterrolebinding:

**kubectl apply -f ~/STW-CaaSPv4/labs/nfs-client-storageclass/auth**

You should see the following were created (not necessarily in this order):

serviceaccount “nfs-client-provisioner”  
clusterrole “nfs-client-provisioner-runner”  
clusterrolebinding “run-nfs-client-provisioner”

6. Enter the following command to view the service accounts:

**kubectl get serviceaccounts**

You should see nfs-client-provisioner listed.

7. Enter the following command to view the cluster roles:

**kubectl get clusterroles**

You should see nfs-client-provisioner-runner listed.

### Task 3: Define a Storage Class

1. On the CAAS101-management VM, in the text editor of your choice, open the file:

**~/STW-CaaSPv4/labs/nfs-client-storageclass/deploy/class.yaml**

```
apiVersion: v1beta1
kind: ServiceAccount
metadata:
  name: nfs-client-provisioner
```

2. In the text editor of your choice, open the file:

**~/STW-CaaSPv4/labs/nfs-client-storageclass/deploy/deployment.yaml**

```
kind: Deployment
apiVersion: extensions/v1beta1
metadata:
  name: nfs-client-provisioner
spec:
  replicas: 1
  strategy:
    type: Recreate
  template:
    metadata:
      labels:
        app: nfs-client-provisioner
    spec:
      containers:
        - name: nfs-client-provisioner
          image: rmt.example.com:5000/quay.io/external_storage/nfs-client-provisioner:latest
          volumeMounts:
            - name: nfs-client-root
              mountPath: /persistentvolumes
      env:
        - name: PROVISIONER_NAME
          value: rmt.example.com
        - name: NFS_SERVER
          value: 192.168.110.2
        - name: NFS_PATH
          value: /export/volumes
  serviceAccount: nfs-client-provisioner
  serviceAccountName: nfs-client-provisioner
  volumes:
    - name: nfs-client-root
      nfs:
        server: 192.168.110.2
        path: /export/volumes
```

3. Enter the following command to create the storage class and provisioner:

**kubectl apply -f ~/STW-CaaSPv4/labs/nfs-client-storageclass/deploy**

You should see the following were created (not necessarily in this order):

storageclass “managed-nfs-storage”  
deployment “nfs-client-provisioner”

4. Enter the following command to view the service accounts:

**kubectl get storageclasses**

You should see **managed-nfs-storage** listed.

5. Enter the following command to view the cluster roles:

**kubectl get deployments**

You should see **nfs-client-provisioner** listed.

## Task 4: Create a Persistent Volume Claim and Pod that Uses It

6. On the CAAS101-management VM, in the text editor of your choice, open the file:

**~/STW-CaaSPv4/labs/nfs-client-storageclass/test/test-claim.yaml**

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: test-claim
  annotations:
    volume.beta.kubernetes.io/storage-class: "managed-nfs-storage"
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 1Mi
```

7. In the text editor of your choice, open the file:

**~/STW-CaaSPv4/labs/nfs-client-storageclass/test/test-pod.yaml**

```
kind: Pod
apiVersion: v1
metadata:
  name: test-pod
spec:
  containers:
    - name: test-pod
      image: rmt.example.com:5000/gcr.io/google_containers/busybox:1.24
      command:
        - "/bin/sh"
      args:
        - "-c"
        - "touch /mnt/SUCCESS && exit 0 || exit 1"
      volumeMounts:
        - name: nfs-pvc
          mountPath: "/mnt"
    restartPolicy: "Never"
  volumes:
    - name: nfs-pvc
      persistentVolumeClaim:
        claimName: test-claim
```

8. Enter the following command to create the volume claim and the pod:

**kubectl apply -f ~/STW-CaaSPv4/labs/nfs-client-storageclass/test**

You should see the following were created (not necessarily in this order):

persistentvolumeclaim “test-claim”  
pod “test-pod”

9. Enter the following command to view the persistent volume claims:

### **kubectl get pvc**

You should see **test-claim** listed.

10. Enter the following command to view the details about the volume claim:

### **kubectl describe pvc test-claim**

Under the **Events** section you should see that the provisioner successfully provisioned the volume followed by the name of the volume. Record the name of the volume here:

**NFS\_VOLUME\_NAME=** \_\_\_\_\_

This will be in the directory name you will see on the NFS server

11. Enter the following command to view the pods:

### **kubectl get pods**

You should note see **test-pod** listed because it was defined to deploy, run a command to write a file to the volume and then exit.

12. On the **SMT server**, enter the following command to view the directory that was created for the volume:

### **ls -l /export/volumes/**

You should see a directory named:

**default-test-claim-NFS\_VOLUME\_NAME**

13. Enter the following command:

### **ls -l /export/volumes/default-test-claim-NFS\_VOLUME\_NAME**

You should see a file named: **SUCCESS**

## **Task 5: Create a Persistent Volume Claim and Pod that Uses It**

14. On the **CAAS101-management VM**, enter the following command to remove the volume claim:

### **kubectl delete pvc test-claim**

You should see that the persistentvolumeclaim was deleted.

15. Enter the following command to view the persistent volume claims:

### **kubectl get pvc**

You should no longer see **test-claim** listed.

16. On the **SMT server**, enter the following command to view the directory that was created for the volume:

### **ls -l /export/volumes/**

You should see that the directory for the volume has been renamed to:

**archived-default-test-claim-NFS\_VOLUME\_NAME**

If desired, you could apply the manifests in the test directory again to see another directory created.

## Task 6: Remove the Objects from the Cluster

1. On the CAAS101-management VM, in the first terminal, enter the following commands to delete all of the objects:

```
kubectl delete -f ~/STW-CaaSPv4/labs/nfs-client-storageclass/test
kubectl delete -f ~/STW-CaaSPv4/labs/nfs-client-storageclass/deploy
kubectl delete -f ~/STW-CaaSPv4/labs/nfs-client-storageclass/auth
```

You should see that the objects were deleted.

---

### **Summary:**

In this exercise, you created manifests for a service account, a role, a role binding, a storageclass, a persistent volume claim, a pod to that attached to the volume and writes data to a file in the volume. You then verified that the volume was being created and the file in it was created on the NFS server.

(End of Exercise)

## 6 Deploying Advanced Workload

---

### Description:

How to deploy a workload from an existing Docker Image

### 6- 1 Run Heimdall as a standalone container

---

### Description:

In this exercise, you configure and run Heimdall as standalone container

---

#### Task 1: Create a place for Heimdall to store it's files

1. On the Workstation, enter the following commands in a terminal session to create the ~/docker/heimdall folder:

```
sudo mkdir -p /docker/heimdall  
sudo chmod 777 /docker -R  
cd /docker/heimdall
```

The directory and file should now exist.

#### Task 2: Run Heimdall as a container

1. On the Workstation, enter the following commands in a terminal session to create

```
sudo docker run -p 80:80 -v /docker/heimdall:/config -e  
TZ=Europe/London rmt.example.com:5000/heimdall
```

This tells docker to run the heimdall container with the following options:

**Map port 80 in the container to 80 on the host**

**mount ~/docker/heimdall in the container in /config**

**Set the timezone within the container**

2. Notice that the first time it runs it created a default installation

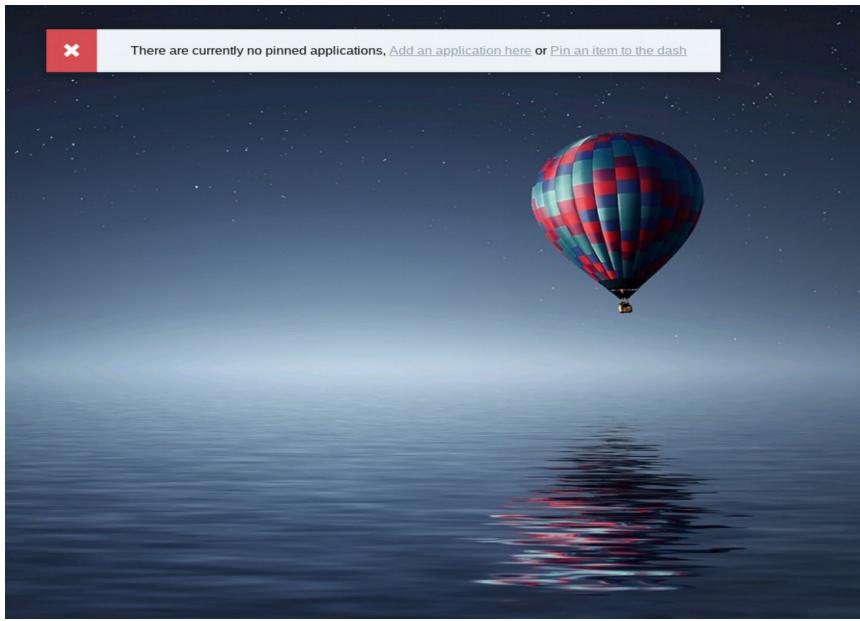
```
-----  
      ()  
     /--\ | | /--\ | |  
    \__\ \_\ \_\ \_\ /  
  
Brought to you by linuxserver.io  
We gratefully accept donations at:  
https://www.linuxserver.io/donate/  
-----  
GID/UID  
-----  
  
User uid: 911  
User gid: 911  
-----  
  
[cont-init.d] 10-adduser: exited 0.  
[cont-init.d] 20-config: executing...  
[cont-init.d] 20-config: exited 0.  
[cont-init.d] 30-keygen: executing...  
generating self-signed keys in /config/keys, you can rep  
Generating a RSA private key  
.....+++++  
.....  
writing new private key to '/config/keys/cert.key'  
----  
[cont-init.d] 30-keygen: exited 0.  
[cont-init.d] 50-config: executing...  
New container detected, installing Heimdall  
Creating app key. This may take a while on slower system  
Application key set successfully.  
Setting permissions  
[cont-init.d] 50-config: exited 0.  
[cont-init.d] 99-custom-scripts: executing...  
[custom-init] no custom scripts found exiting...  
[cont-init.d] 99-custom-scripts: exited 0.  
[cont-init.d] done.  
[services.d] starting services  
[services.d] done.  
■
```

### Task 3:On the Workstation launch the Chrome Browser

- 1 On the Workstation **launch the Chrome Browser**

<http://127.0.0.1>

- 2 You should see a default screen with no Apps Defined



## Task 4: Create an App in Helm

1. Click on 'Add an application here'

**Application Name:** SUSE

**URL:** <http://www.suse.com>

**Select:** Pinned

Add application

PINNED    SAVE    CANCEL

Application name *	Application Type *	Colour *
SUSE	None	Hex colour
URL	Tags (Optional)	
<a href="http://www.suse.com">http://www.suse.com</a>		 Upload a file
Preview		
 		
		SAVE    CANCEL

2. Click Save to Save the App

## Task 5: Stop and restart container

1. On the Workstation, go back to the terminal session running the container and press the following Keys to stop the container instance

**<ctrl>C**

This will terminate the container instance

2. Restart the Docker Container

```
sudo docker run -p 80:80 -v ~/docker/heimdall:/config -e  
TZ=Europe/London smt.example.com:5000/heimdall
```

Notice How it start also immediately because it's using the configuration from the last time we launched the Application

## Task 6: Verify with Browser

1. On the Workstation launch the Chrome Browser

<http://12.0.0.1>

2. You should see the App you previously defined

## Task 7: Stop Container Instance

1. On the Workstation, go back to the terminal session running the container and press the following Keys to stop the container instance

**<ctrl>C**

This will terminate the container instance

## Task 8: Try and run multiple versions of Heimdall

1. On the Workstation, go back to the terminal session running the container and press the following Keys to stop the container instance

```
sudo docker run -d -p 80:80 -v ~/docker/heimdall:/config -e  
TZ=Europe/London smt.example.com:5000/heimdall
```

**\* notice we added the -d command to tell docker to run as a Daemon rather than in the foreground**

1. Try and run a second instance of the Heimdall

```
sudo docker run -d -p 80:80 -v ~/docker/heimdall:/config -e TZ=Europe/London smt.example.com:5000/heimdall
```

**Notice is fails because the port is already in use**

---

#### **Summary:**

In this exercise, you launched a container from the command line to see how it behaved

(End of Exercise)

---

## **6- 2 Launch Heimdall**

---

#### **Description:**

In this exercise, we will Heimdall without dedicated storage

---

### **Task 1: Create the Manifest for the Persistent Volume**

1. On the Workstation, in the text editor of your choice, open the file:

```
~/STW-CaaSPv4/labs/heimdall/app-health/heimdall-deployment_ns.yaml
```

Enter the following in the file:

## SUSE U Advanced - SUSE CaaS Platform

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: heimdall-shared-deployment
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: heimdall-shared
    spec:
      containers:
        - name: heimdall-pod
          image: linuxserver/heimdall
          ports:
            - containerPort: 80
```

2. On the Workstation, in the text editor of your choice, open the file:

**`~/STW-CaaSPv4/labs/heimdall/app-health/heimdall-service.yaml`**

Enter the following in the file:

```
apiVersion: v1
kind: Service
metadata:
  name: heimdall-shared-service
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 33000
  selector:
    app: heimdall-shared
```

## Task 2: Deploy Heimdall

1. From a Terminal prompt on the Management workstation

**`kubectl apply -f ~/STW-CaaSPv4/labs/heimdall/app-health/heimdall-deployment_ns.yaml`**

**`kubectl apply -f ~/STW-CaaSPv4/labs/heimdall/app-health/heimdall-service.yaml`**

## Task 3: View the Deployment

- From Kubernetes Dashboard select Pods and then click on the Pod that was just deployed
- Click on the Log button so we can see what is happening on that Pod

The screenshot shows the Kubernetes Dashboard interface. In the top navigation bar, the URL is `localhost:8001/api/v1/namespaces/kube-system/services/https:kubernetes-dashboard:/proxy/#/pod/default/heimdall-shared-deployment-c8df77bf4-jz2s2`. Below the URL, there's a search bar and several action buttons: '+ CREATE', 'LOGS' (which is highlighted with a red box), 'EDIT', and 'DELETE'. The main content area is titled 'Details' and contains the following information for the pod 'heimdall-shared-deployment-c8df77bf4-jz2s2':
 

Name:	heimdall-shared-deployment-c8df77bf4-jz2s2
Namespace:	default
Labels:	app: heimdall-shared, pod-template-hash: 748933690
Annotations:	container.apparmor.security.beta.kubernetes.io/heimdall-pod: runtime/default, kubernetes.io/psp: suse.caasp.psp.unprivileged, seccomp.security.alpha.kubernetes.io/pod: docker/default
Creation Time:	2019-06-05T06:56 UTC
Status:	Running
QoS Class:	BestEffort

- Look at the logs and notice how it is creating a new installation

The screenshot shows the 'Logs' page of the Kubernetes Dashboard. The left sidebar lists various cluster components like Namespaces, Nodes, and Deployments. The main area displays the logs for the 'heimdall-pod' in the 'heimdall-shared-deployment-c8df77bf4-jz2s2' deployment. The logs output by the Heimdall keygen process are as follows:
 

```

    ()_
    | |
    | | / _\ | | / _\ 
    | | \_ \ | | | () |
    |_ | _/_ |_ \_/
    -----
    Brought to you by linuxserver.io
    We gratefully accept donations at:
    https://www.linuxserver.io/donate/
    -----
    GID/UID
    -----
    User uid: 911
    User gid: 911
    -----
    [cont-init.d] 10-adduser: exited 0.
    [cont-init.d] 20-config: executing...
    [cont-init.d] 20-config: exited 0.
    [cont-init.d] 30-keygen: executing...
    generating self-signed keys in /config/keys, you can replace these with your own keys if required
    Generating a RSA private key
    ..+++++
    .....+++++
    writing new private key to '/config/keys/cert.key'
    -----
    [cont-init.d] 30-keygen: exited 0.
    [cont-init.d] 50-config: executing...
    New container detected, installing Heimdall
    Creating app key. This may take a while on slower systems
    -----
    
```

SUSE U Advanced - SUSE CaaS Platform

4. Let that deployment continue to run until it is completely finished (Refresh the Browser to check Status)

## Task 4: View Heimdall in a Browser

1. On the **Workstation** open the **Chrome browser** and go to  
**worker.example.com:33000**
  
2. Notice how ever time it is deployed it has to rebuild the installation
3. Stop the deployment by either killing it on the Kubernetes Dashboard from the command line by typing in the following command

**kubectl delete -f ~/STW-CaaSPv4/labs/heimdall/app-health/heimdall-deployment\_ns.yaml**

## Task 4: Re-deploy Heimdall

1. Repeat all of the steps in **Task 2 and Task 3**
2. Notice how ever time it is deployed it has to rebuild the installation

---

### Summary:

In this exercise, you created manifests for a Heimdall. We then deployed it and watched it build the default configuration. We learned that every time we deployed the app it had to create the default installation.

(End of Exercise)

---

## 6- 3 Configure Heimdall to use NFS Persistent Storage

---

### Description:

In this exercise, you configure a persistent volume on an NFS server. You then launch a pod that uses the persistent volume

---

## Task 1: Create the Manifest for the Persistent Volume

1. On the **Workstation**, in the text editor of your choice, open the file:  
**~/STW-CaaSPv4/labs/heimdall/app-shared/heimdall-pv-shared.yaml**

Enter the following in the file:

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: heimdall-shared
spec:
  capacity:
    storage: 200Mi
  accessModes:
    - ReadWriteMany
  nfs:
    server: 192.168.110.2
    path: "/export/heimdall-shared"
```

## Task 2: Create the Manifest for the Persistent Volume Claim

1. In the text editor of your choice, open the file:

**~/STW-CaaSPv4/labs/heimdall/app-shared/heimdall-pvc-shared.yaml**

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: heimdall-shared-claim
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 200Mi
```

## Task 3: Create the Manifest for the Service so we can access the deployment

1. In the text editor of your choice, open the file:
2. **~/STW-CaaSPv4/labs/heimdall/app-shared/heimdall-service.yaml**

```
apiVersion: v1
kind: Service
metadata:
  name: heimdall-shared-service
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 30000
  selector:
    app: heimdall-shared
```

#### Task 4: Create the Manifests for the Web Frontend

1. In the text editor of your choice, create/open the file:
2. **~/STW-CaaSPv4/labs/heimdall/app-shared/heimdall-deployment.yaml**

```

apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: heimdall-shared-deployment
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: heimdall-shared
    spec:
      containers:
        - name: heimdall-pod
          image: linuxserver/heimdall
          volumeMounts:
            - name: heimdall-pvc-shared
              mountPath: "/config"
      volumes:
        - name: heimdall-pvc-shared
          persistentVolumeClaim:
            claimName: heimdall-shared-claim

```

## Task 5: Deploy the Objects

1. To deploy the volumes/pods/service, open a terminal and enter the following command:

**kubectl apply -f ~/STW-CaaSPv4/labs/heimdall/app-shared**

You should see the following were created (not necessarily in this order):

deployment.extensions "heimdall-shared-deployment" created  
 persistentvolume "heimdall-shared" created  
 persistentvolumeclaim "heimdall-shared-claim" created  
 service "heimdall-shared-service" created

2. Enter the following command to view the deployments:

**kubectl get deployments**

You should see the **heimdall-shared-deployment** deployments listed.

3. Enter the following command to view the persistent volumes:

**kubectl get pv**

You should see the persistent volume **heimdall-shared** listed.

4. Enter the following command to view the persistent volumes:

**kubectl get pvc**

You should see the persistent volume claim **heimdall-shared-claim** listed.

## Task 6: Test the Persistent Data

1. On the **management workstation**, open a web browser and point to:

**http://worker10.example.com:30000**

Notice we have a Heimdall session pre-populated with 2 Applications

## Task 7: Remove the Objects from the Cluster

1. In the first terminal, enter the following commands to delete all of the objects:

**kubectl delete -f ~/STW-CaaSPv4/labs/heimdall/app-shared**

You should see that the objects were deleted.

---

### Summary:

In this exercise, you created manifests for a persistent volume, a persistent volume claim, a pod to that attached to the volume and writes data to an index.html file in the volume, and a web server that attaches to the volume and displays the index.html file. You then verified that the index.html file was being updated by looking at the file both on the NFS volume and the web server.

(End of Exercise)